

Heat Transfer

Contact Information:

Bianca Palade (ACES) email: bianca.palade@stud.fim.upb.ro,
palade.c.bianca@gmail.com

Introduction

Physical simulations can be among the most computationally challenging problems to solve. Fundamentally, there is often a trade-off between accuracy and computational complexity. As a result, computer simulations have become more and more important in recent years, thanks in large part to the increased accuracy possible because of the parallel computing revolution. Since many physical simulations can be parallelized quite easily, we will look at a very simple simulation model in this example.

In this paper, it simulated the heat transfer in a Microfluidic Device. The program receives an image with sensor model, computes heat transfer for a specific number of steps and make a short video with heat animation.

Objective

- Implementation of the serial implementation of the algorithm
- Implementation of the algorithm to check for the correctness of the proposed method
- Implementation of the parallel implementation of the algorithm and optimize the results for the proposed paradigm

Background

The implementation of this application requires:

- C++ skills
- basic OpenCV skills
- Math
- Image Processing

Resources

- https://computing.llnl.gov/tutorials/parallel_comp/
- Barbara Chapman et al. , "Using OpenMP: Portable Shared Memory Parallel Programming (Scientific and Engineering Computation)" , The MIT Press, 2007.
- Gregory V. Wilson, Paul Lu, "Parallel programming using C++" , The MIT Press, 1996.
- <https://www.coursera.org/learn/parallelism-ia/home/welcome>

Mathematical model

The law of heat conduction, also known as Fourier's law, states that the rate of heat transfer through a material is proportional to the negative gradient in the temperature and to the area, at right angles to that gradient, through which the heat flows.

The differential form of Fourier's law of thermal conduction shows that the local heat flux density q , is equal to the product of thermal conductivity, k , and the negative local temperature gradient, $-\nabla T$. The heat flux density is the amount of energy that flows through a unit area per unit time.

$$q = -k\nabla T$$

where (including the SI units)

q is the local heat flux density, $\text{W}\cdot\text{m}^{-2}$

k is the material's conductivity, $\text{W}\cdot\text{m}^{-1}\cdot\text{K}^{-1}$,

∇T is the temperature gradient, $\text{K}\cdot\text{m}^{-1}$.

In my heat conduction model, we will compute the heat in every pixel of the picture. For each point of mesh we will solve numerical heat as a sum of old temperature and product of material's conductivity and the difference between current temperature and the temperatures of its neighbor (right, left, top, bottom).

$$q = T_{old} + k(T_{right} + T_{left} + T_{top} + T_{bottom} - 4 \cdot T_{old})$$

Algorithm description

Import Input Data

In this project, the algorithm will take the geometric model data from a bmp file. Where the black contour will represent the geometric shape of the heat resistance with max temperature T_{MAX} , and the white spaces, will represent the environment with the minimum temperature T_{MIN} .

```

pix      = cv::imread("mems.bmp", 1);
pix_out = pix;

if (pix.empty())
    std::cout << "Fail to open!" << std::endl;
else
    std::cout << "Image opened fine!" << std::endl;

picture_in = (float*)malloc(DIM * DIM * sizeof(float));
picture_out = (float*)malloc(DIM * DIM * sizeof(float));
picture_ct = (float*)malloc(DIM * DIM * sizeof(float));

for (int i = 0; i < DIM; i++) {
    for (int j = 0; j < DIM; j++) {
        picture_ct[DIM * i + j] = picture_in[DIM * i + j] =
            (pix.at<cv::Vec3b>(i, j)[0] == 0) ? T_MAX : T_MIN;
    }
}

```

To create the short video animation we used the VideoWriter class of the OpenCV 3.0 library.

```
VideoWriter video("outcpp.avi", cv::VideoWriter::fourcc('M', 'J', 'P', 'G'), 10, Size(DIM, DIM));
```

The video is MJPEG compress, with a frame rate of 10 fps and DIM * DIM resolution.

The algorithm is structured in 2 main parts: heat transfer computing and conversion of numerical values in RGB space.

Heat Transfer Computation

Before calculating the heat transfer, I need to save the initial data in a constant image, which will be kept throughout the simulation so as not to lose the initial conditions. This will be used every step of the way to calculate heat transfer.

```
void copy_values_from_previous_step(float* in, float* ct) {
    for (int i = 0; i < DIM * DIM; i++) {
        if (ct[i] != T_MIN)
            in[i] = ct[i];
    }
}
```

Computation of the output temperatures is based on previous temperatures. In function update_values will be compute and save the temperatures. Then we will swap the input array with the output array.

```
void update_values(float* out, float* in) {
    long long left, right, top, bottom;

    for (long long i = 0; i < DIM * DIM; i++) {
        left = (i % DIM == 0) ? i : i - 1;
        right = (i % DIM == DIM - 1) ? i : i + 1;
        top = (i < DIM) ? i : i - DIM;
        bottom = (i >= DIM * (DIM - 1)) ? i : i + DIM;

        out[i] = in[i] + 0.1 * (in[top] + in[left] + in[right] + in[bottom] - in[i] * 4);
    }
}
```

Conversion from Float to RGB

Next step is conversion of floating-point temperatures values in RGB space. We will assume that temperatures values will be in range [0.0001,1] where 1 will be the fraction between maximum temperature and maximum temperature and 0.0001 the ratio between minimum temperature and maximum temperature. In color heat transfer I will consider that we work in HLS space and we will convert HLS values to RGB. In our color space we consider:

- lightness = temperature,
- saturation = 1,
- hue = (180 + (int)(360.0f * lightness)) % 360

```
unsigned char value(float n1, float n2, int hue)
{
    hue = (hue > 360) ? (hue - 360) : ((hue < 0) ? hue + 360 : hue);
    if (hue < 60)
        return (unsigned char)(255 * (n1 + (n2 - n1) * hue / 60));
    if (hue < 180)
        return (unsigned char)(255 * n2);
    if (hue < 240)
        return (unsigned char)(255 * (n1 + (n2 - n1) * (240 - hue) / 60));
    return (unsigned char)(255 * n1);
}
```

```
void float_to_color(Mat ptr_out, float* out){
    // Find pixel position
    for (int i = 0; i < DIM; i++)
        for(int j = 0; j < DIM; j++){
            float lightness = out[DIM * i + j];
            float saturation = 1;
            int hue = (180 + (int)(360.0f * lightness)) % 360;
            float m1, m2;

            m2 = (lightness <= 0.5f) ? lightness * (1 + saturation) :
                lightness + saturation - lightness * saturation;

            m1 = 2 * lightness - m2;

            ptr_out.at<cv::Vec3b>(i,j)[2] = value(m1, m2, hue + 120);
            ptr_out.at<cv::Vec3b>(i,j)[1] = value(m1, m2, hue);
            ptr_out.at<cv::Vec3b>(i,j)[0] = value(m1, m2, hue - 120);
        }
}
```

Main Function

```
bool dest = true;
for (int i = 0; i < 100; i++) {
    for (int i = 0; i < CONDUCTIVITY; i++) {
        if (dest) {
            copy_values_from_previous_step(picture_in, picture_ct);
            update_values(picture_out, picture_in);
        }
        else {
            copy_values_from_previous_step(picture_out, picture_ct);
            update_values(picture_in, picture_out);
        }
        dest = !dest;
    }
    float_to_color(pix_out, picture_out);
    video.write(pix_out);
}
video.release();
```

Profiling for Sequential Program

Profiling for Sequential Program

I used Intel VTune for profiling.

Collection and Platform Info

This section provides information about this collection, including result set size and collection platform data.

Application Command Line: C:\Users\bianca.palade\source\repos\Heat_Transfer_Sequential\x64\Debug\Heat_Transfer_Sequential.exe

Environment Variables:

Operating System: Microsoft Windows 10

Computer Name: DESKTOP-TFFSL40

Result Size: 3 MB

Collection start time: 14:36:40 29/11/2021 UTC

Collection stop time: 14:36:47 29/11/2021 UTC

Collector Type: User-mode sampling and tracing

Finalization mode: Fast. If the number of collected samples exceeds the threshold, this mode limits the number of processed samples.

CPU

Name: Intel(R) Processor code named Kabylake ULX

Frequency: 2.7 GHz

Logical CPU Count: 4

Cache Allocation Technology

Level 2 capability: not detected

Level 3 capability: not detected


Hotspots by CPU Utilization

Elapsed Time: 45.827 s

Total Thread : 1

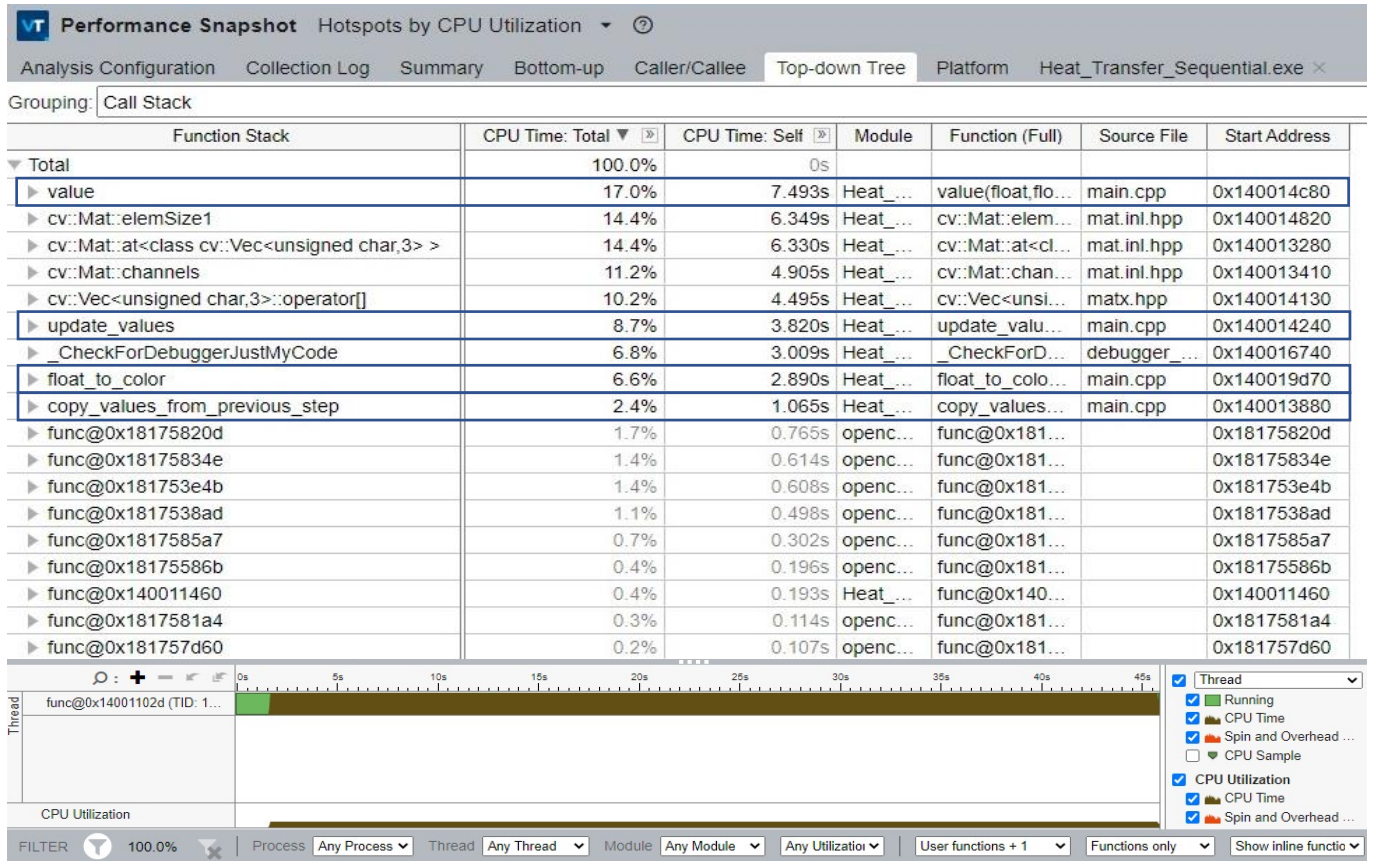
Paused Time : 0 s

Effective CPU Utilization Histogram



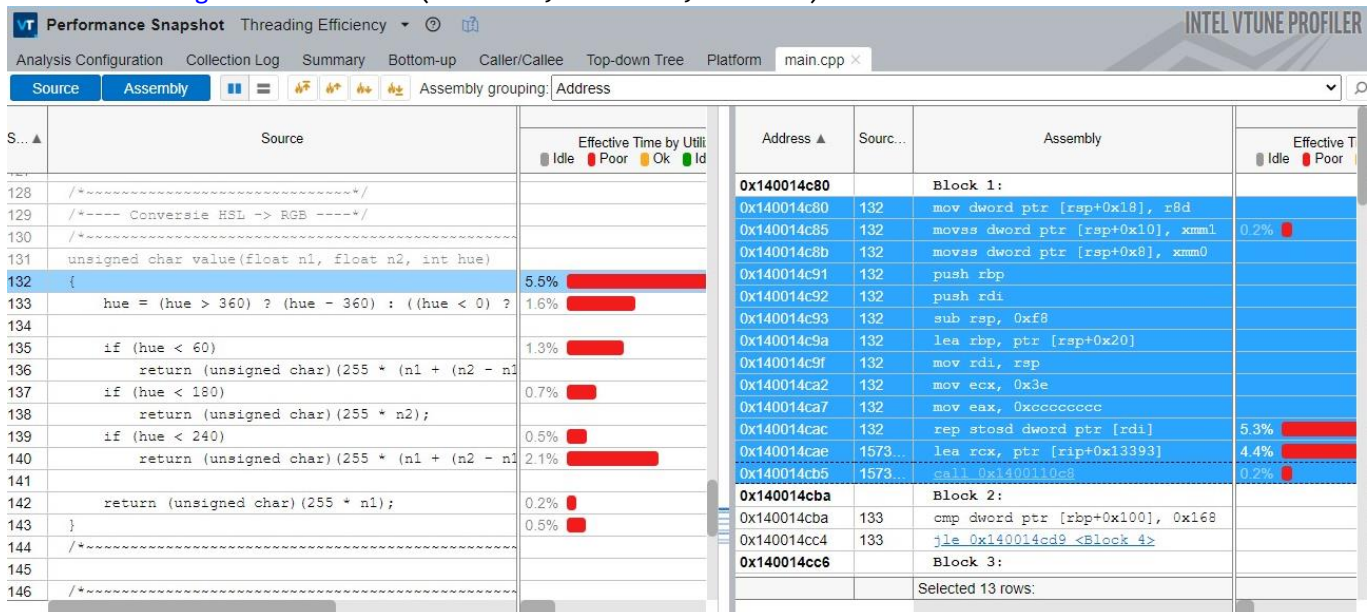


Top-down Tree



Hotspots Functions (source + assembly)

1. unsigned char value (float n1, float n2, int hue)



2. void update_values(float* out, float* in)

Performance Snapshot Threading Efficiency				INTEL VTUNE PROFILER			
Analysis Configuration Collection Log Summary Bottom-up Caller/Callee Top-down Tree Platform main.cpp x main.cpp x							
Source Assembly				Assembly grouping: Address			
Source	Effective Time by Ut	Address	Source Line	Assembly	Effective Ti		
108		0x14001430	122	mov rcx, qword ptr [rbp+0x28]	0.0%		
109		0x1400143fa	122	addss xmm1, dword ptr [rax+rcx*4]			
110		0x140014403	122	mov rax, qword ptr [rbp+0x188]	0.5%		
111		0x14001440a	122	mov rcx, qword ptr [rbp+0x68]			
112		0x14001440e	122	addss xmm1, dword ptr [rax+rcx*4]			
113		0x140014413	122	mov rax, qword ptr [rbp+0x188]	0.6%		
114		0x14001441a	122	mov rcx, qword ptr [rbp+0x88]			
115		0x140014421	122	movss xmm2, dword ptr [rax+rcx*4]			
116		0x140014426	122	mulss xmm2, dword ptr [rip+0xac8e]			
117		0x14001442e	122	subss xmm1, xmm2			
118		0x140014432	122	cvtss2sd xmm1, xmm1	0.2%		
119		0x140014436	122	movsd xmm2, qword ptr [rip+0xac62]	1.1%		
120		0x14001443e	122	mulsd xmm2, xmm1			
121		0x140014442	122	movaps xmm1, xmm2	0.7%		
122		0x140014445	122	addsd xmm0, xmm1			
123		0x140014449	122	cvtss2ss xmm0, xmm0	0.2%		
124		0x14001444d	122	mov rax, qword ptr [rbp+0x180]	1.1%		
125		0x140014454	122	mov rcx, qword ptr [rbp+0x88]			
126		0x14001445b	122	movss dword ptr [rax+rcx*4], xmm0			

Parallel Algorithm

To parallelize the algorithm, I used OpenMP. I tried to improve the execution time by parallelizing the hotspots functions

1. `void float_to_color(Mat ptr_out, float* out)`

The value function `value()` cannot be parallelized efficiently, because it only checks conditions and

```
void float_to_color(Mat ptr_out, float* out)
{
    int i, j;
    #pragma omp parallel for num_threads(MAX_THREADS) private(i,j)
    for (i = 0; i < DIM; i++)
        for (j = 0; j < DIM; j++) {
            float lightness = out[DIM * i + j];
            float saturation = 1;
            int hue = (180 + (int)(360.0f * lightness)) % 360;
            float m1, m2;

            m2 = (lightness <= 0.5f) ? lightness * (1 + saturation) : lightness +
            saturation - lightness * saturation;
            m1 = 2 * lightness - m2;

            ptr_out.at<cv::Vec3b>(i, j)[2] = value(m1, m2, hue + 120);
            ptr_out.at<cv::Vec3b>(i, j)[1] = value(m1, m2, hue);
            ptr_out.at<cv::Vec3b>(i, j)[0] = value(m1, m2, hue - 120);}}
```

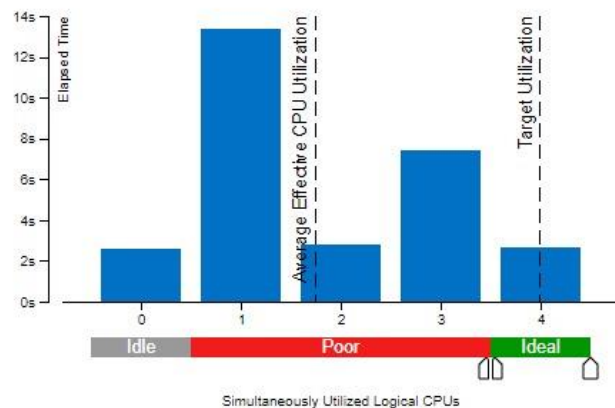
does not contain any repetitive instructions, but the `float_to_color()` function, can be parallelized because it calls the `value()` function for `DIM * DIM * 3` times.


Profiling

Hotspots by CPU Utilization

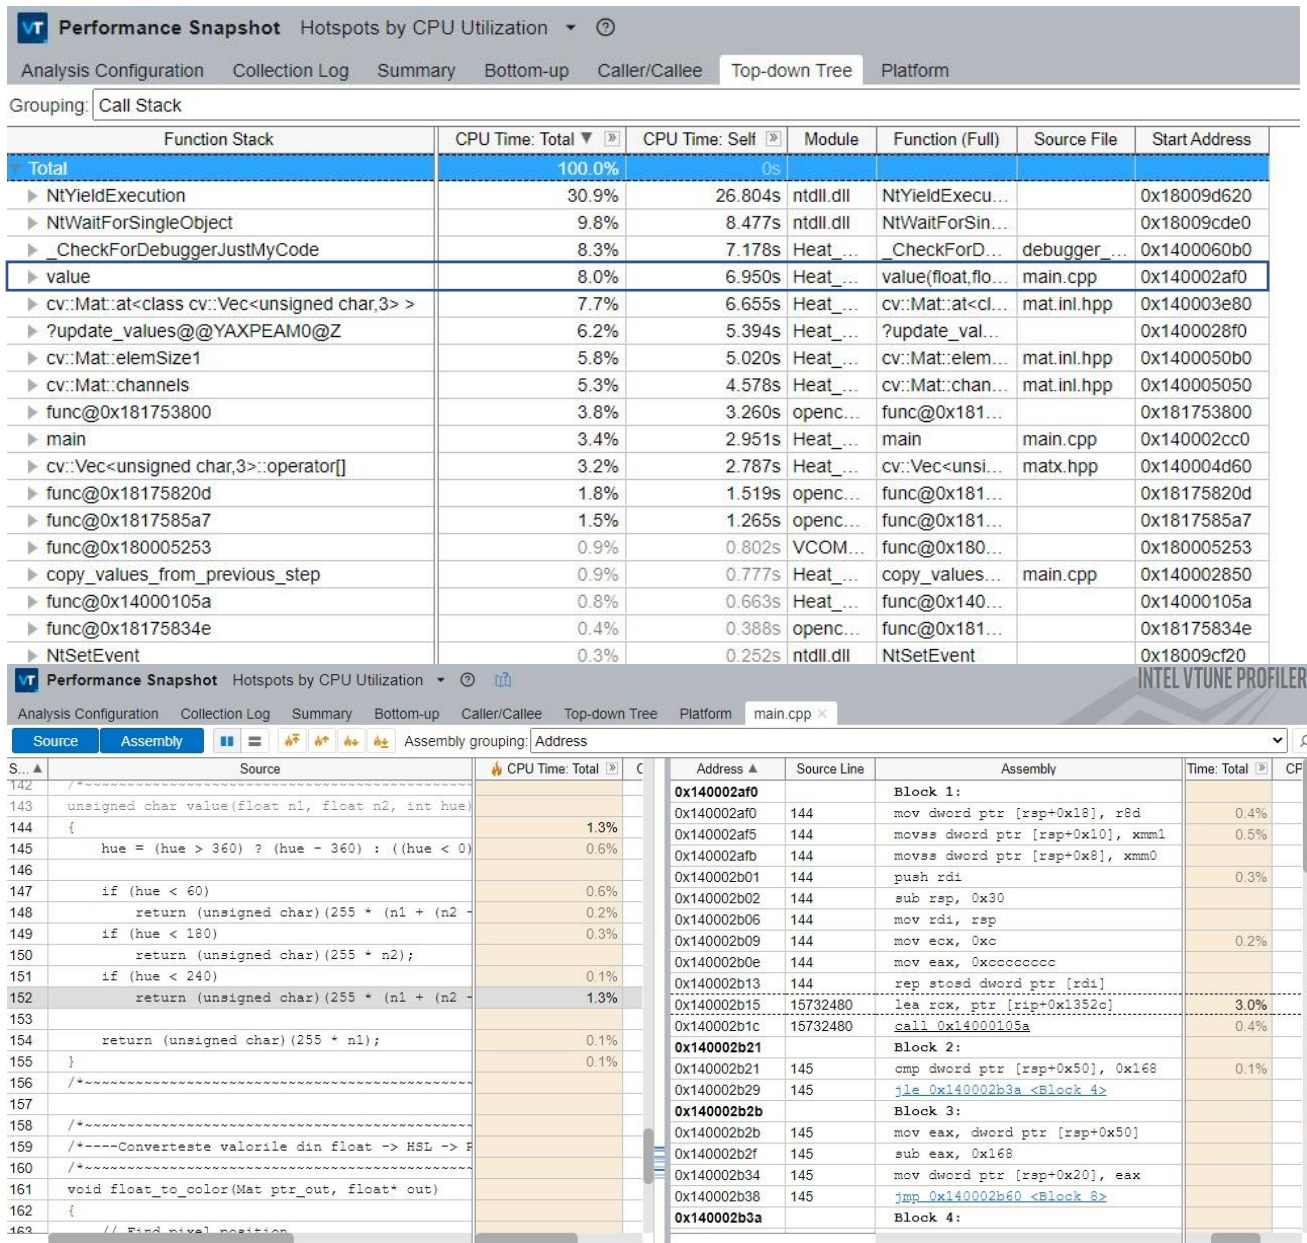
Effective CPU Utilization Histogram

Elapsed Time: 28.852 s
Total Thread : 4
Paused Time : 0 s





Hotspots Functions (source + assembly)



2. `void update_values(float* out, float* in)`

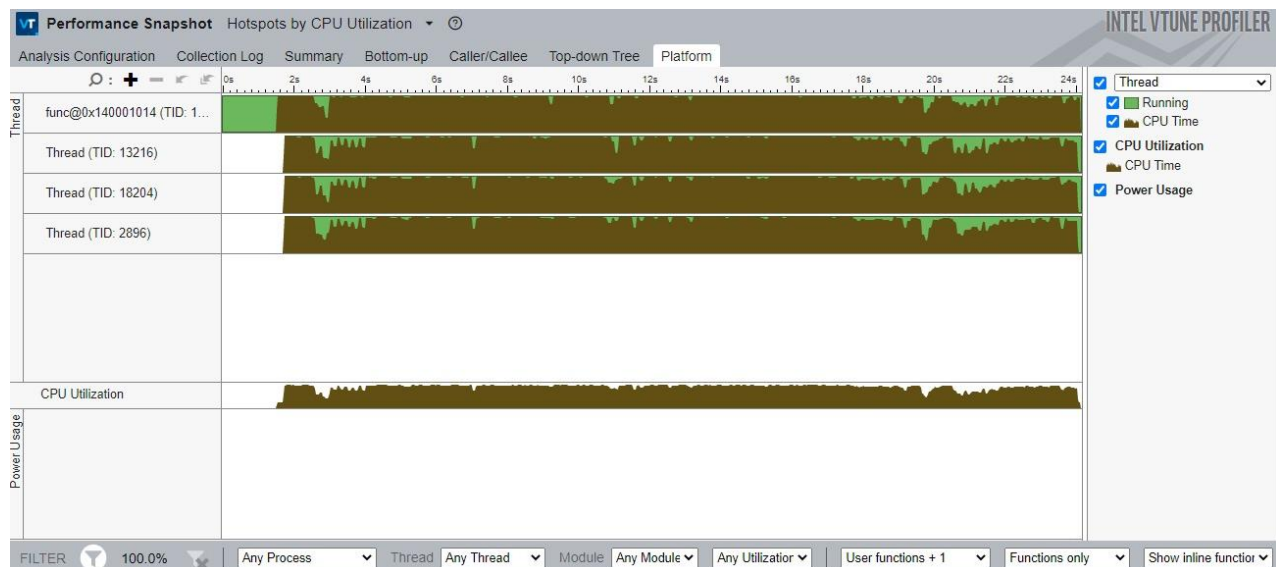

The second most time consuming is the `update_values()` function. To parallelize it, I divided the vector `out` into the maximum number of threads.

```
void update_values(float* out, float* in) {
    #pragma omp parallel
    for (long long i = omp_get_thread_num() * DIM * DIM / MAX_THREADS; i <
        omp_get_thread_num() * DIM * DIM / MAX_THREADS + DIM * DIM / MAX_THREADS; ++i) {
        long long left = (i % DIM == 0) ? i : i - 1;
        long long right = (i % DIM == DIM - 1) ? i : i + 1;
        long long top = (i < DIM) ? i : i - DIM;
        long long bottom = (i >= DIM * (DIM - 1)) ? i : i + DIM;
        out[i] = in[i] + 0.1f * (in[top] + in[left] + in[right] + in[bottom] - in[i]
            * 4);
    }
}
```

Hotspots by CPU Utilization

Elapsed Time: 23.587 s
Total Thread : 4
Paused Time : 0 s

Effective CPU Utilization Histogram



Initially, I tried to use `#pragma omp parallel for ordered`, but the program run much slower than the sequential version of the function, because when OpenMP distributes the work among threads there is a lot of administration/synchronisation going on to ensure the values in my vector are not corrupted somehow.

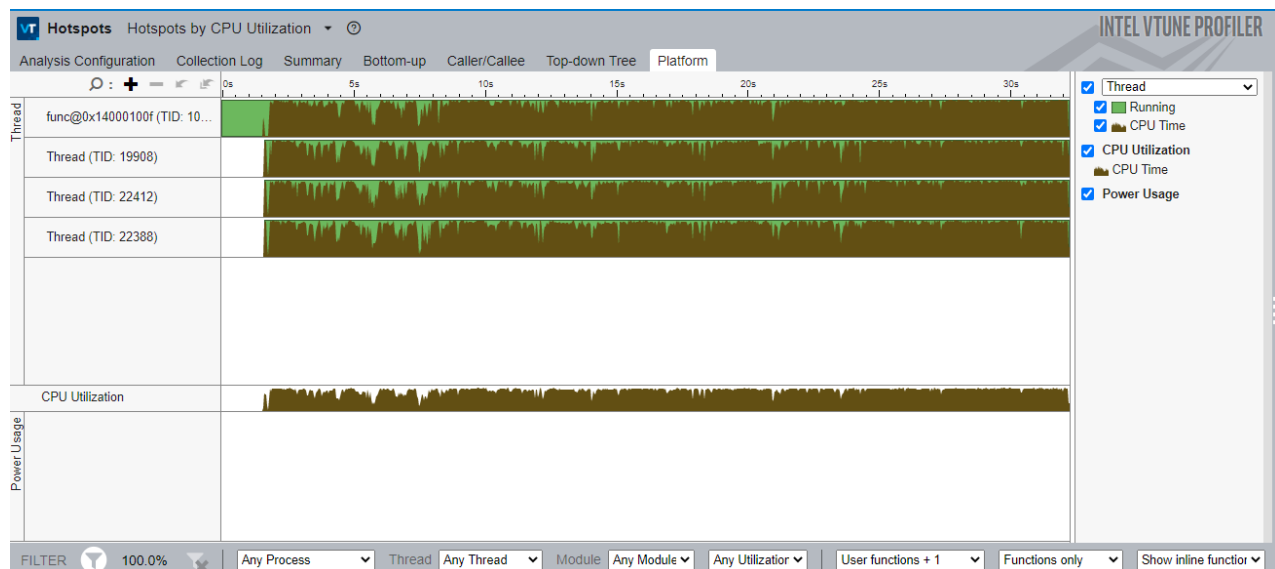

```
#pragma omp parallel for ordered
for (long long i = 0; i < DIM * DIM; i++) {
    left = (i % DIM == 0) ? i : i - 1;
    right = (i % DIM == DIM - 1) ? i : i + 1;
    top = (i < DIM) ? i : i - DIM;
    bottom = (i >= DIM * (DIM - 1)) ? i : i + DIM;

    out[i] = in[i] + 0.1 * (in[top] + in[left] + in[right] + in[bottom] - in[i]
* 4);}
```

Hotspots by CPU Utilization

Effective CPU Utilization Histogram

Elapsed Time: 32.215 s
Total Thread : 4
Paused Time : 0 s



3. `void copy_values_from_previous_step(float* in, float* ct)`

```
void copy_values_from_previous_step(float* in, float* ct) {
    #pragma omp parallel for num_threads(MAX_THREADS)
    for (int i = 0; i < DIM * DIM; i++) {
        if (ct[i] != T_MIN)
            in[i] = ct[i];
    }
}
```

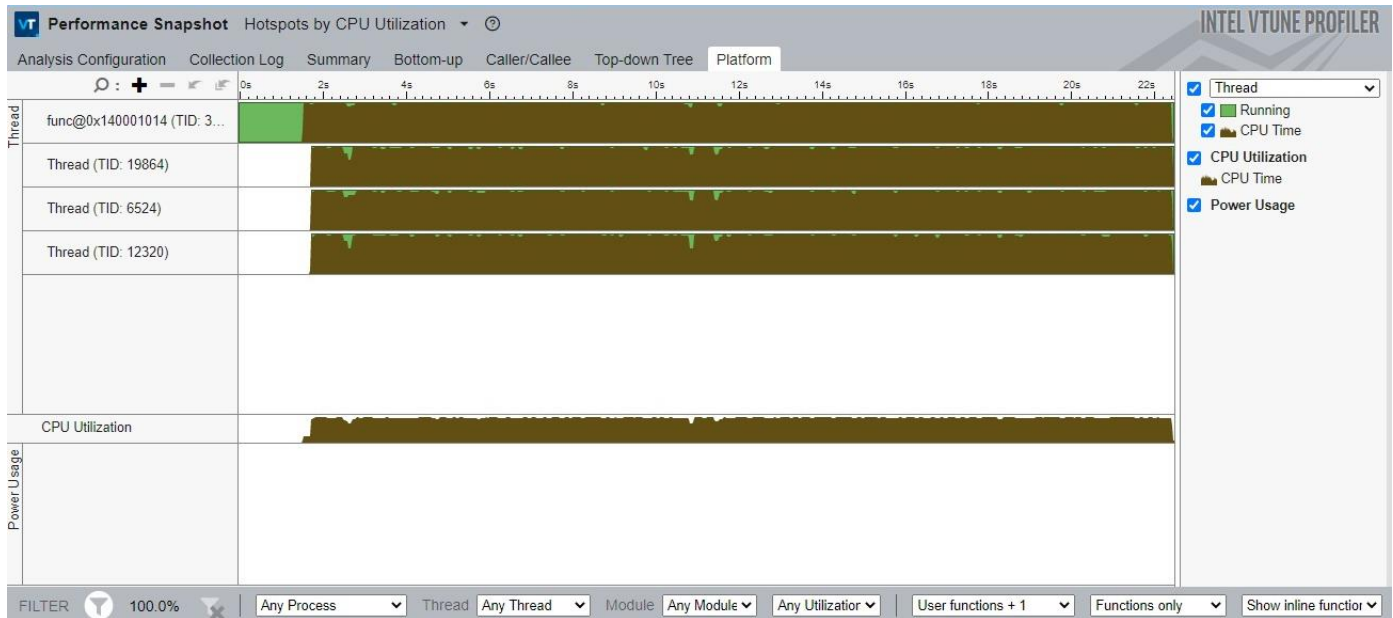

Hotspots by CPU Utilization

Effective CPU Utilization Histogram

Elapsed Time: 22.452 s

Total Thread : 4

Paused Time : 0 s

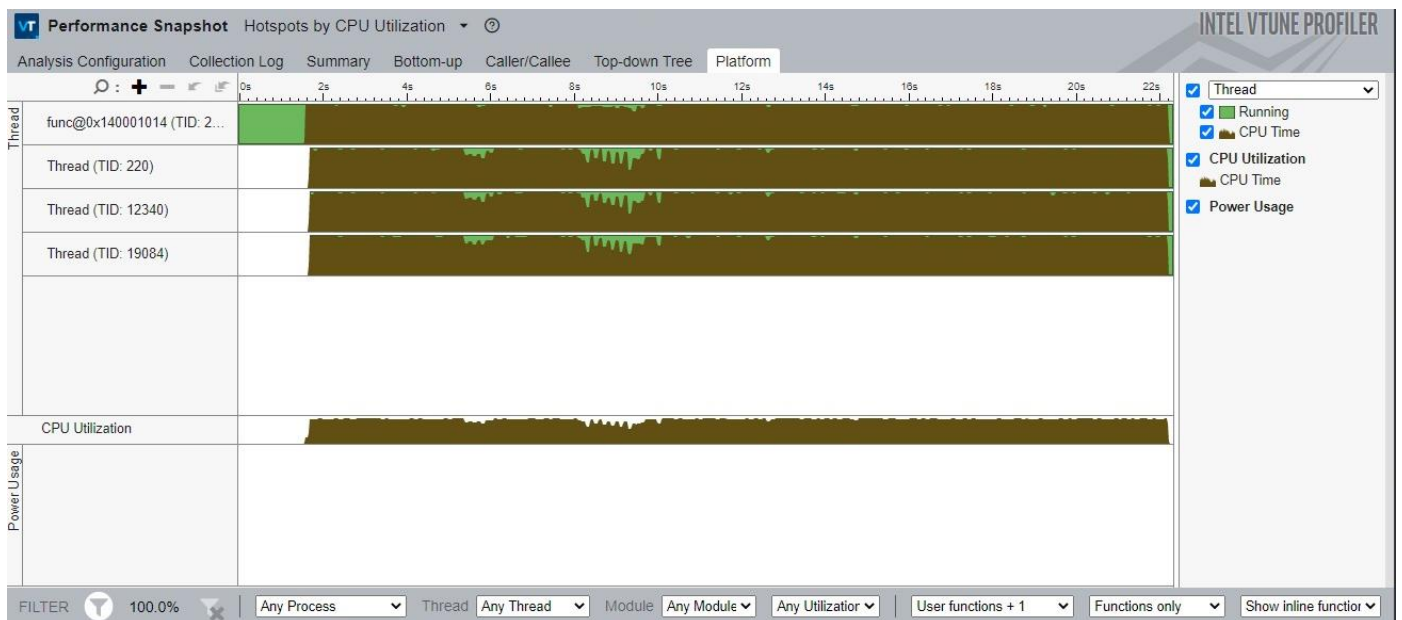
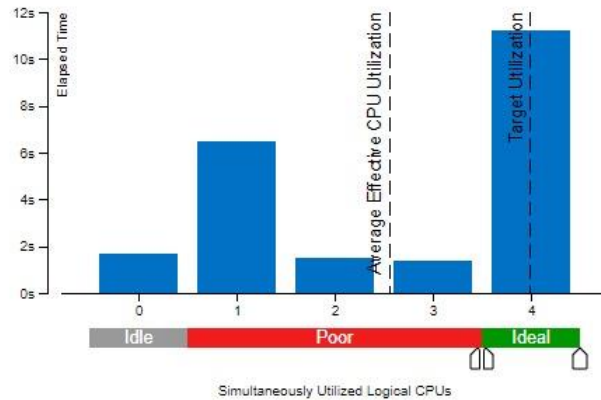


```
#pragma omp parallel for num_threads(MAX_THREADS) private(i,j)
for (i = 0; i < DIM; i++) {
    for (j = 0; j < DIM; j++) {
        picture_in[DIM * i + j] = (pix.at<cv::Vec3b>(i, j)[0] == 0) ? T_MAX : T_MIN;
        picture_ct[DIM * i + j] = picture_in[DIM * i + j];
    }
}
```

4. from int main I parallelized image read

Hotspots by CPU Utilization
Effective CPU Utilization Histogram

Elapsed Time: 22.303 s
Total Thread : 4
Paused Time : 0 s



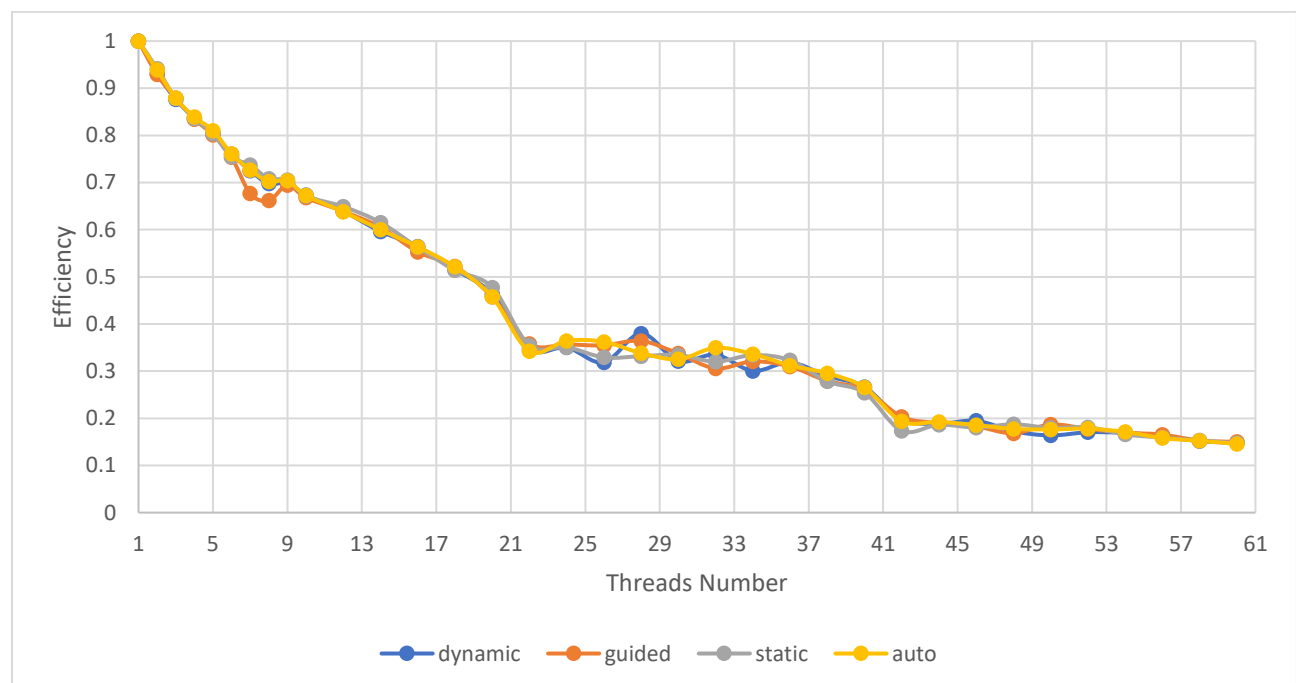
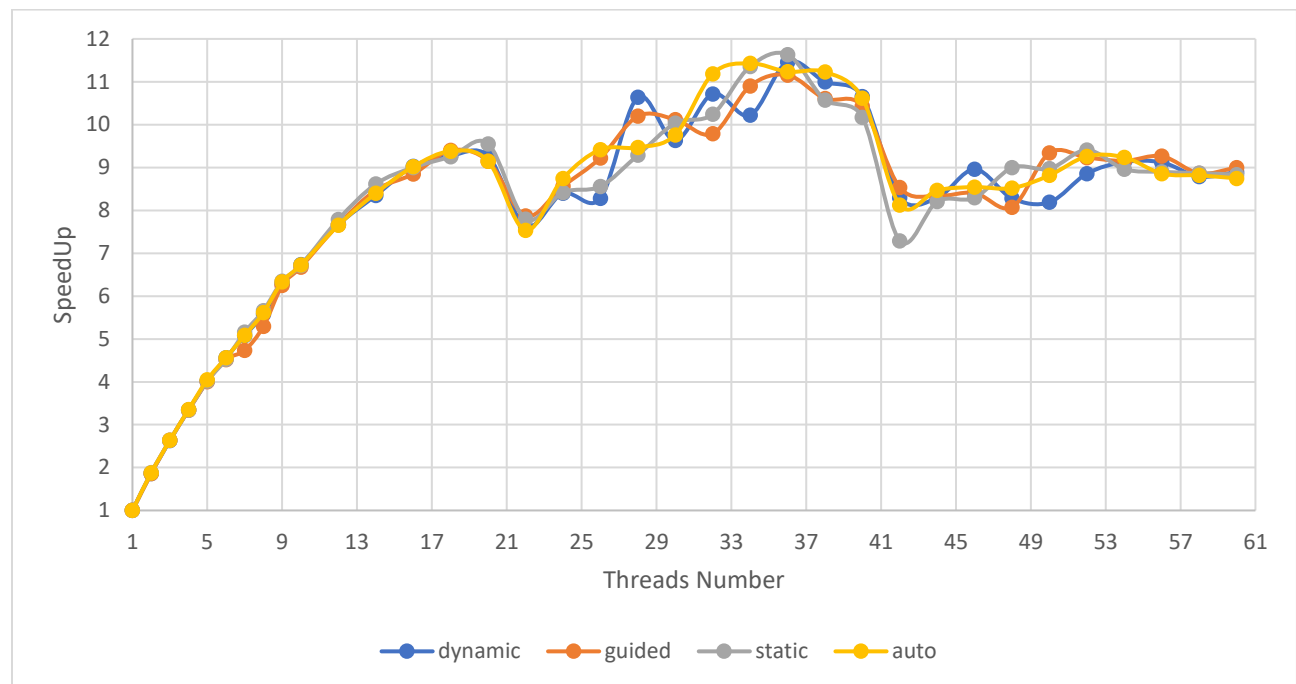
Hotspots Functions

Performance Snapshot Hotspots by CPU Utilization						
Analysis Configuration Collection Log Summary Bottom-up Caller/Callee Top-down Tree Platform						
Grouping: Call Stack						
Function Stack	CPU Time: Total	CPU Time: Self	Module	Function (Full)	Source File	Start Address
▼ Total	100.0%	0s				
► NtYieldExecution	25.2%	19.783s	ntdll.dll	NtYieldExecu...		0x18009d620
► main	13.7%	10.772s	Heat_...	main	main.cpp	0x140009570
► _CheckForDebuggerJustMyCode	10.1%	7.929s	Heat_...	_CheckForD...	debugger_...	0x1400060b0
► cv::Mat::at<class cv::Vec<unsigned char,3> >	9.9%	7.800s	Heat_...	cv::Mat::at<cl...	mat.inl.hpp	0x140003e80
► value	8.6%	6.785s	Heat_...	value(float, flo...	main.cpp	0x1400093a0
► cv::Mat::channels	6.3%	4.929s	Heat_...	cv::Mat::chan...	mat.inl.hpp	0x140003e30
► cv::Vec<unsigned char,3>::operator[]	4.9%	3.855s	Heat_...	cv::Vec<unsi...	matx.hpp	0x140004aa0
► cv::Mat::elemSize1	4.6%	3.637s	Heat_...	cv::Mat::elem...	mat.inl.hpp	0x140004090
► func@0x181753e4b	2.4%	1.901s	openc...	func@0x181...		0x181753e4b
► func@0x1817538f5	2.1%	1.649s	openc...	func@0x181...		0x1817538f5
► func@0x18175820d	1.9%	1.470s	openc...	func@0x181...		0x18175820d
► func@0x180027390	1.4%	1.098s	KERN...	func@0x180...		0x180027390
► func@0x14000105a	1.3%	1.060s	Heat_...	func@0x140...		0x14000105a
► omp_get_thread_num	1.3%	1.007s	VCOM...	omp_get_thr...		0x1800059f0
► func@0x1817585a7	1.1%	0.848s	openc...	func@0x181...		0x1817585a7
► func@0x18175834e	1.0%	0.779s	openc...	func@0x181...		0x18175834e
► func@0x180005254	0.9%	0.676s	VCOM...	func@0x180...		0x180005254
► ?elemSize1@Mat@cv@@@QEBA_KXZ	0.6%	0.496s	Heat_...	?elemSize1...		0x140004090
► TlsGetValue	0.6%	0.491s	KERN...	TlsGetValue		0x180015540
► RtlGetCurrentUmsThread	0.4%	0.304s	ntdll.dll	RtlGetCurren...		0x180058990
► SwitchToThread	0.4%	0.294s	KERN...	SwitchToThr...		0x1800676a0
► func@0x1400011fe	0.3%	0.200s	Heat_...	func@0x140...		0x1400011fe
► GetTickCount	0.2%	0.188s	KERN...	GetTickCount		0x180015640
► func@0x181fa7f30	0.1%	0.110s	openc...	func@0x181f...		0x181fa7f30
► func@0x181fa7cd0	0.1%	0.108s	openc...	func@0x181f...		0x181fa7cd0
► func@0x181757250	0.1%	0.106s	openc...	func@0x181757250		0x181757250
► SwitchToThread	0.1%	0.100s	KERN...	SwitchToThr...		0x18001b400
► func@0x181758552	0.1%	0.093s	openc...	func@0x181...		0x181758552
► func@0x1817558a0	0.1%	0.091s	openc...	func@0x181...		0x1817558a0

Speedup, Efficiency, CPU-time



Running without O3 on HPSL cluster

I ran the program 3 times for each condition, and I did the arithmetic mean of the values





Running with -O3 on HPSL cluster



1. Static, Chunk Size = 4,8,16,32,64 Speedup + Efficiency



2. Guided, Chunk Size = 4,8,16,32,64 Speedup + Efficiency

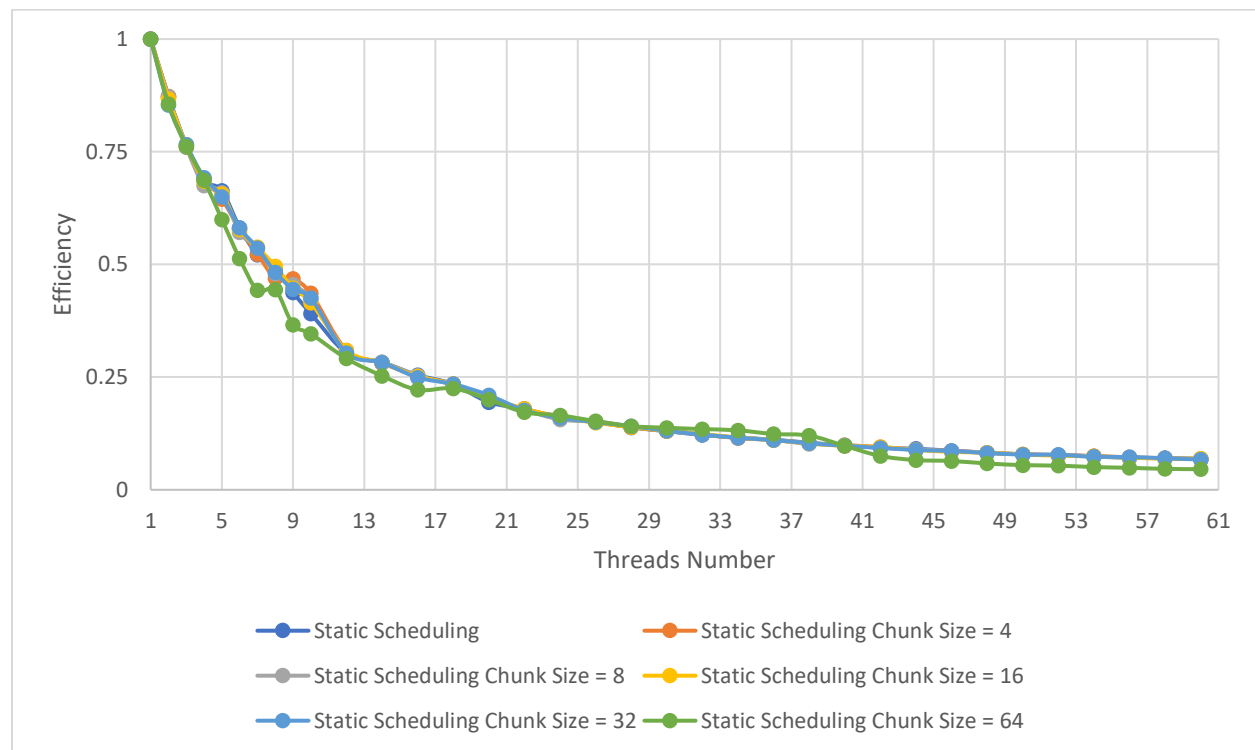



3. Dynamic, Chunk Size = 4,8,16,32,64 Speedup + Efficiency

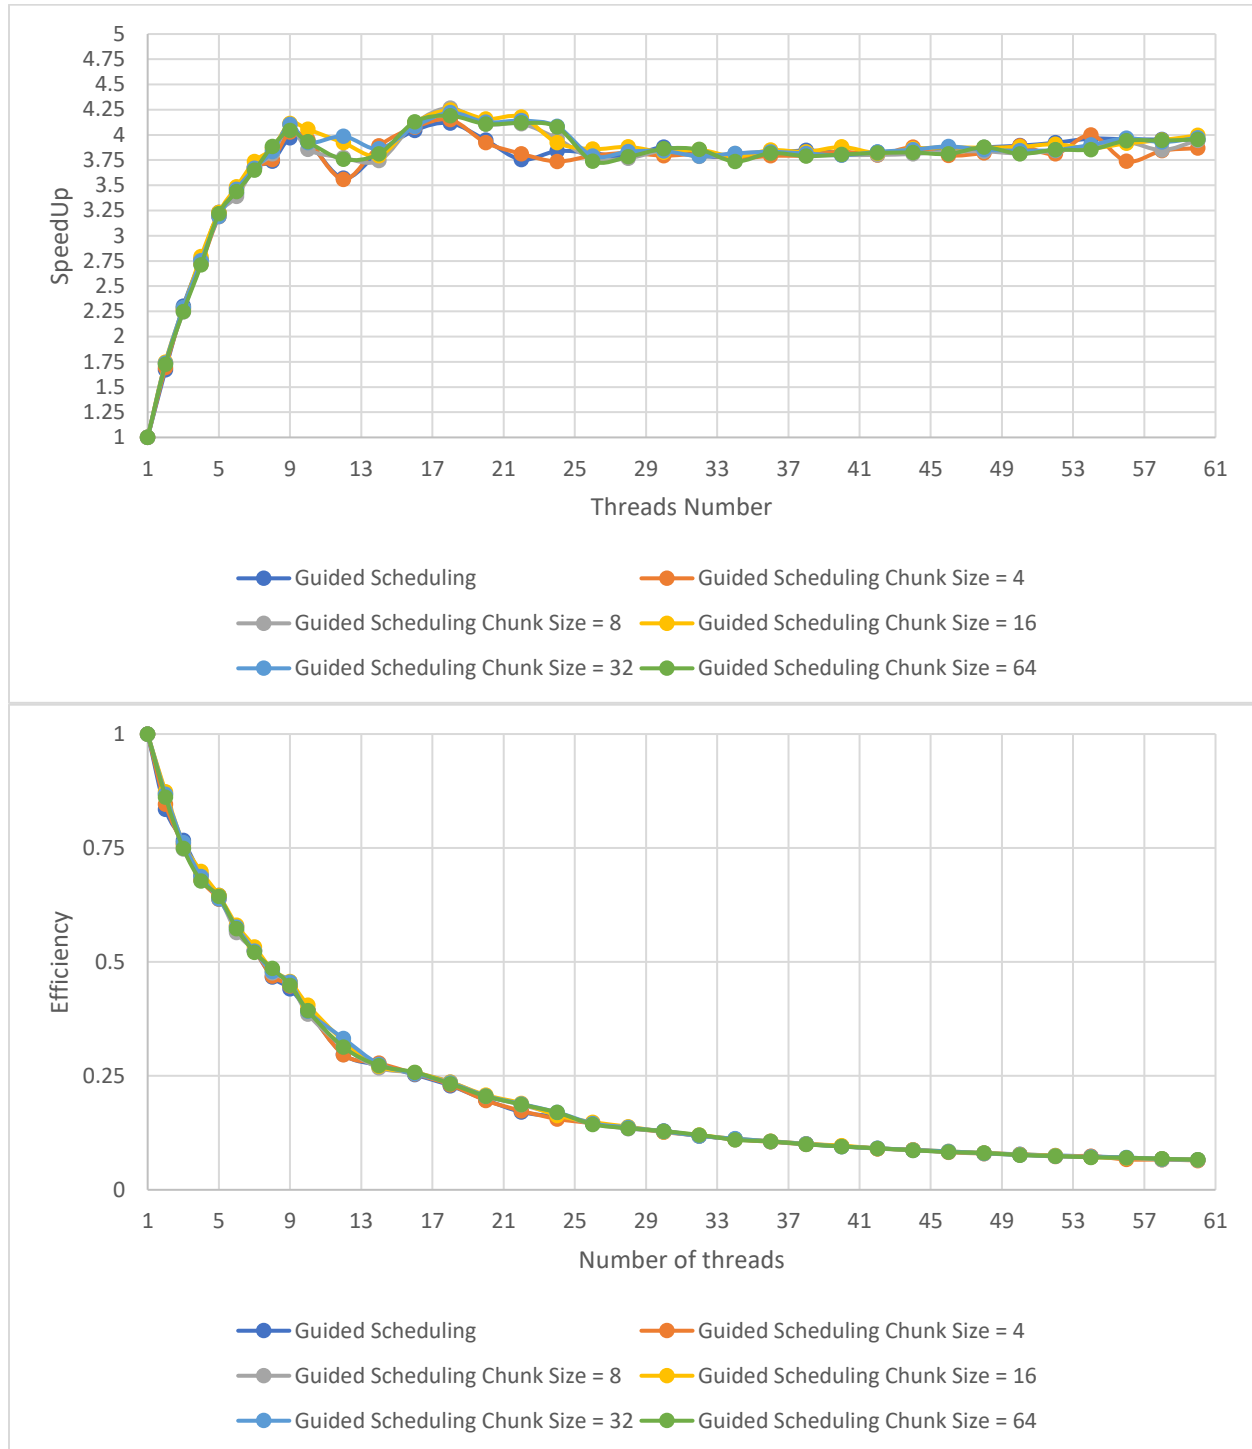


Running with -O3 on IBM cluster


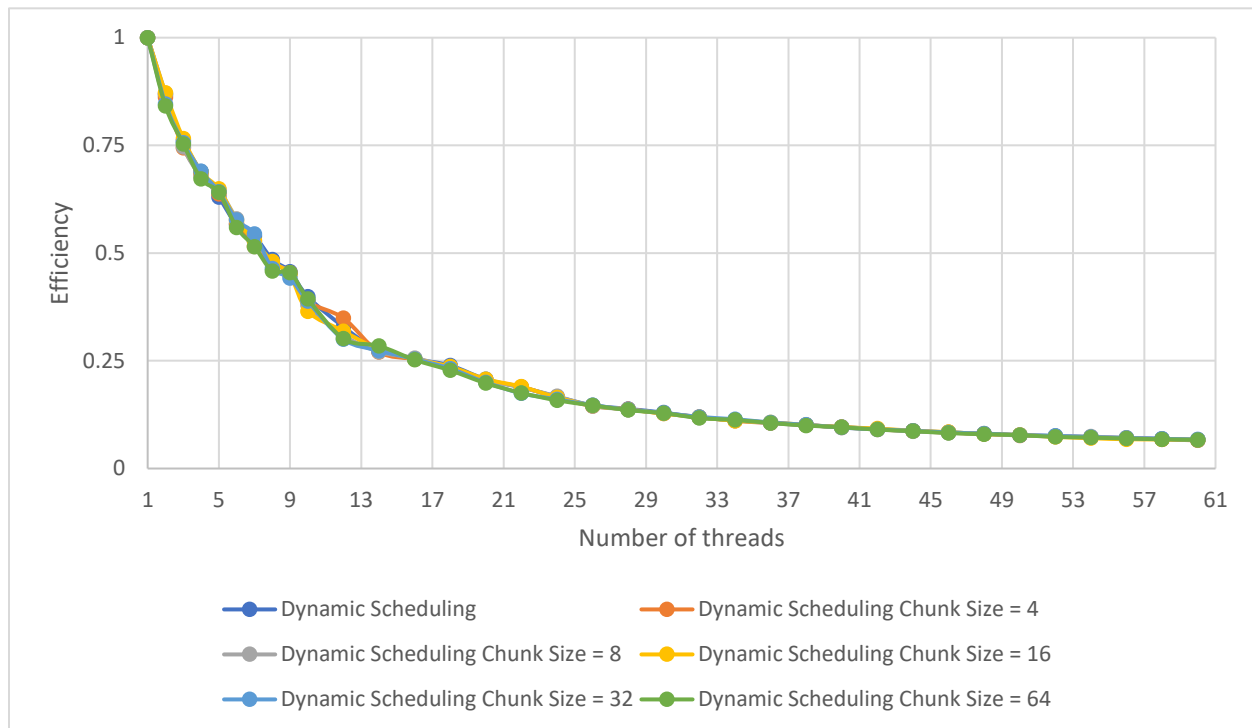
1. Static, Chunk Size = 4,8,16,32,64 Speedup + Efficiency



2. Guided, Chunk Size = 4,8,16,32,64 Speedup + Efficiency



3. Dynamic, Chunk Size = 4,8,16,32,64 Speedup + Efficiency



Conclusion

Speed Up												
Scheduling \ Chunk-size	1		4		8		16		32		64	
	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM
static	4.616	4.221	4.788	4.360	4.530	4.240	4.558	4.223	4.611	4.256	4.571	4.571
	T:36	T: 18	T:38	T: 18	T:38	T: 18	T:38	T: 18	T:38	T: 18	T:38	T: 18
dynamic	4.665	4.303	4.848	4.221	4.722	4.219	4.660	4.258	4.701	4.171	4.744	4.098
	T:38	T:18	T:38	T:18	T:38	T:18	T:36	T:18	T:38	T:18	T:36	T:18
guided	4.671	4.116	4.619	4.149	4.545	4.265	4.700	4.246	4.646	4.220	4.778	4.189
	T: 38	T:18	T:38	T:18	T:34	T:18	T:38	T:18	T:38	T:18	T:38	T:18
auto	HPSL : 4.799											
	T:38											

CPU - time												
Scheduling \ Chunk-size	1		4		8		16		32		64	
	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM
static	12.341	13.940	11,806	13.481	12.528	13.859	12.403	13.874	12.459	13.834	12.449	12.449
	T:36	T:18	T:38	T:18	T:38	T:18	T:38	T:18	T:38	T:18	T:38	T:18
dynamic	12.164	13.813	11.733	13.888	12.033	13.888	12.156	13.811	12.088	13.930	12.022	13.989
	T:38	T:18	T:38	T:18	T:38	T:18	T:36	T:18	T:38	T:18	T:36	T:18
guided	12.254	14.101	12.436	13.894	12.562	13.648	12.051	13.887	12.189	13.992	12.824	13.983
	T:38	T:18	T:38	T:18	T:34	T:18	T:38	T:18	T:38	T:18	T:38	T:18
auto												

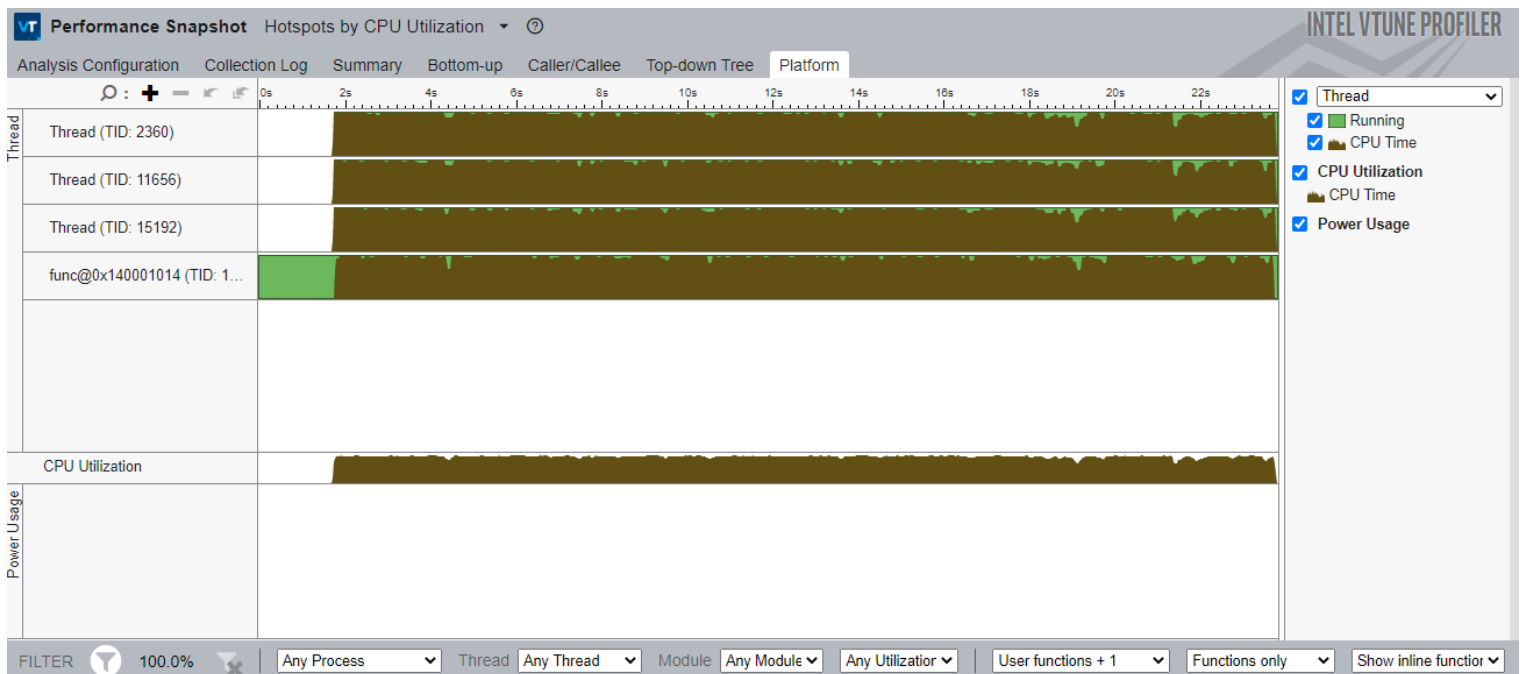

Parallel Algorithm (II)

To parallelize the algorithm, I used OpenMP to split image in number of threads.

Hotspots by CPU Utilization

Effective CPU Utilization Histogram

Elapsed Time: 23.798 s
Total Thread : 4
Paused Time : 0 s



```
int main(int argv, char** argc) {

    . . .

    //Citeste valorile din imagine
    int i, j;
    #pragma omp parallel for num_threads(MAX_THREADS) private(i,j)
    for (i = 0; i < DIM; i++) {
        for (j = 0; j < DIM; j++) {
            picture_ct[DIM * i + j] = picture_in[DIM * i + j] =
            (pix.at<cv::Vec3b>(i, j)[0] == 0) ? T_MAX : T_MIN;
        }
    }

    #pragma omp parallel
    for (int i = 0; i < 100; i++) {
        for (int i = 0; i < CONDUCTIVITY; i++) {
            if (out_pic) {
                copy_values_from_previous_step(picture_in, picture_ct);
                update_values(picture_out, picture_in);
            }
            else {
                copy_values_from_previous_step(picture_out, picture_ct);
                update_values(picture_in, picture_out);
            }
            #pragma omp barrier
            #pragma omp single
            out_pic = !out_pic;
        }

        float_to_color(pix_out, picture_out);
        #pragma omp barrier
        #pragma omp single
        video.write(pix_out);
    }
    video.release();

    . . .

    return 0;
}
```

```
/*~~~~~Copiaza temperatura de la pasul anterior~~~~~*/
/*~~~~~Copiaza temperatura de la pasul anterior~~~~~*/
/*~~~~~Copiaza temperatura de la pasul anterior~~~~~*/
~~~~~*/
void copy_values_from_previous_step(float* in, float* ct) {
    for (long long i = omp_get_thread_num() * DIM * DIM / MAX_THREADS; i <
    omp_get_thread_num() * DIM * DIM / MAX_THREADS + DIM * DIM / MAX_THREADS; ++i) {
        if (ct[i] != T_MIN)
            in[i] = ct[i];
    }
}
/*~~~~~Copiaza temperatura de la pasul anterior~~~~~*/
```

```
/*~~~~~*/
/*----Calculeaza temperatura----*/
/*~~~~~*/
~~~~~*/
void update_values(float* out, float* in) {
    for (long long i = omp_get_thread_num() * DIM * DIM / MAX_THREADS; i <
omp_get_thread_num() * DIM * DIM / MAX_THREADS + DIM * DIM / MAX_THREADS; ++i) {
        long long left = (i % DIM == 0) ? i : i - 1;
        long long right = (i % DIM == DIM - 1) ? i : i + 1;
        long long top = (i < DIM) ? i : i - DIM;
        long long bottom = (i >= DIM * (DIM - 1)) ? i : i + DIM;
        out[i] = in[i] + 0.1f * (in[top] + in[left] + in[right] + in[bottom] - in[i] * 4);
    }
}
/*~~~~~*/
```

```
/*~~~~~*/
/*----Converteste valorile din float -> HSL -> RGB ----*/
/*~~~~~*/
~~~~~*/
void float_to_color(Mat ptr_out, float* out)
{
    // Find pixel position
    for (long long a = omp_get_thread_num() * DIM * DIM / MAX_THREADS; a <
omp_get_thread_num() * DIM * DIM / MAX_THREADS + DIM * DIM / MAX_THREADS; ++a) {
        int i = a / DIM;
        int j = a % DIM;

        float lightness = out[DIM * i + j];
        float saturation = 1;
        int hue = (180 + (int)(360.0f * lightness)) % 360;
        float m1, m2;

        m2 = (lightness <= 0.5f) ? lightness * (1 + saturation) : lightness +
saturation - lightness * saturation;
        m1 = 2 * lightness - m2;

        ptr_out.at<cv::Vec3b>(i, j)[2] = value(m1, m2, hue + 120);
        ptr_out.at<cv::Vec3b>(i, j)[1] = value(m1, m2, hue);
        ptr_out.at<cv::Vec3b>(i, j)[0] = value(m1, m2, hue - 120);
    }
}
/*****/
```

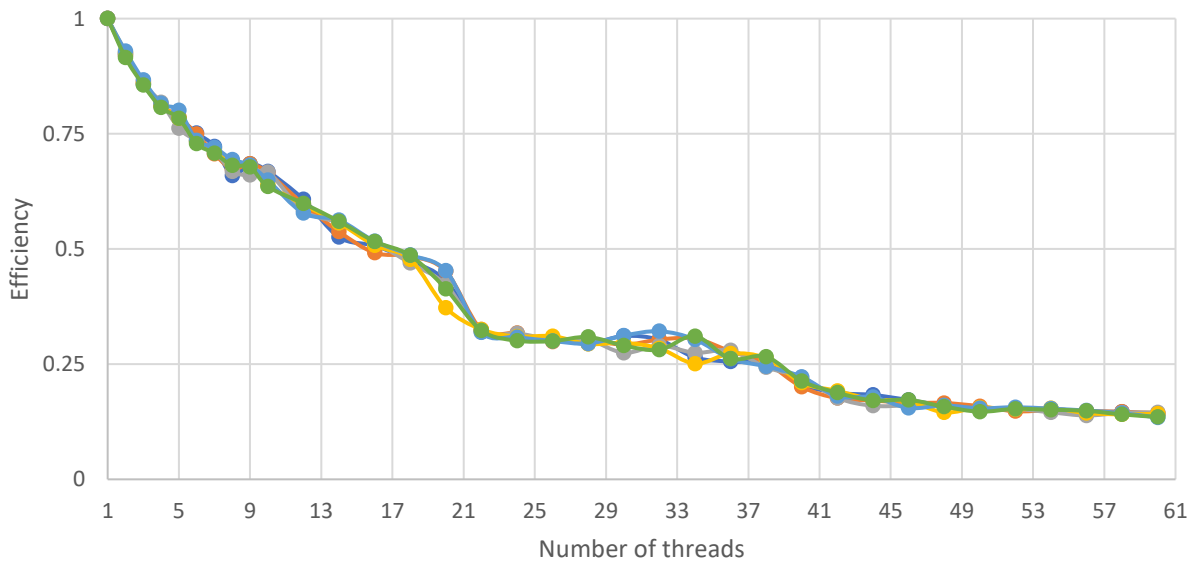
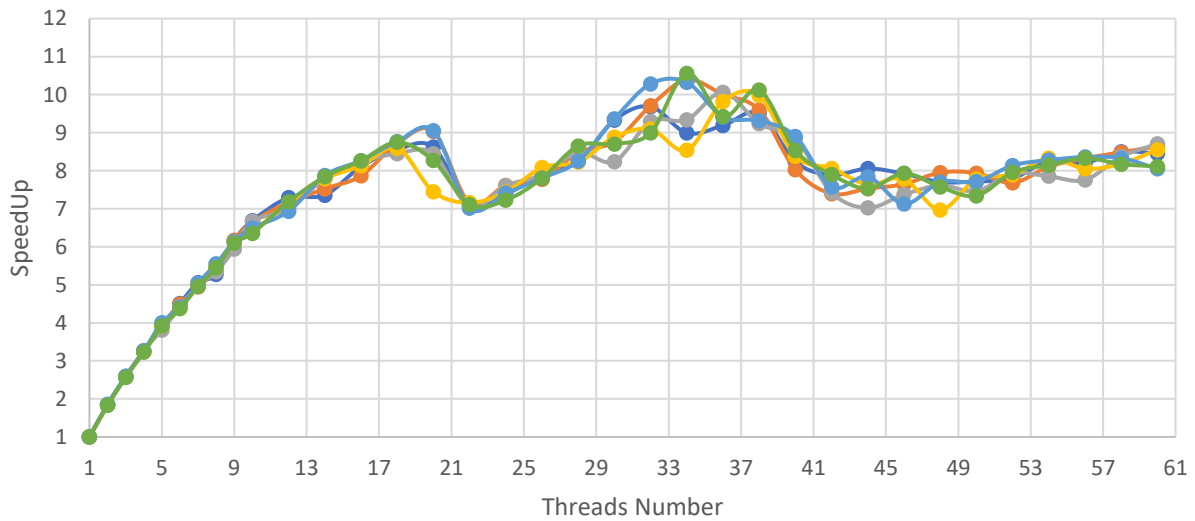
Speedup, Efficiency, CPU-time

Running with -O3 on HPSL cluster

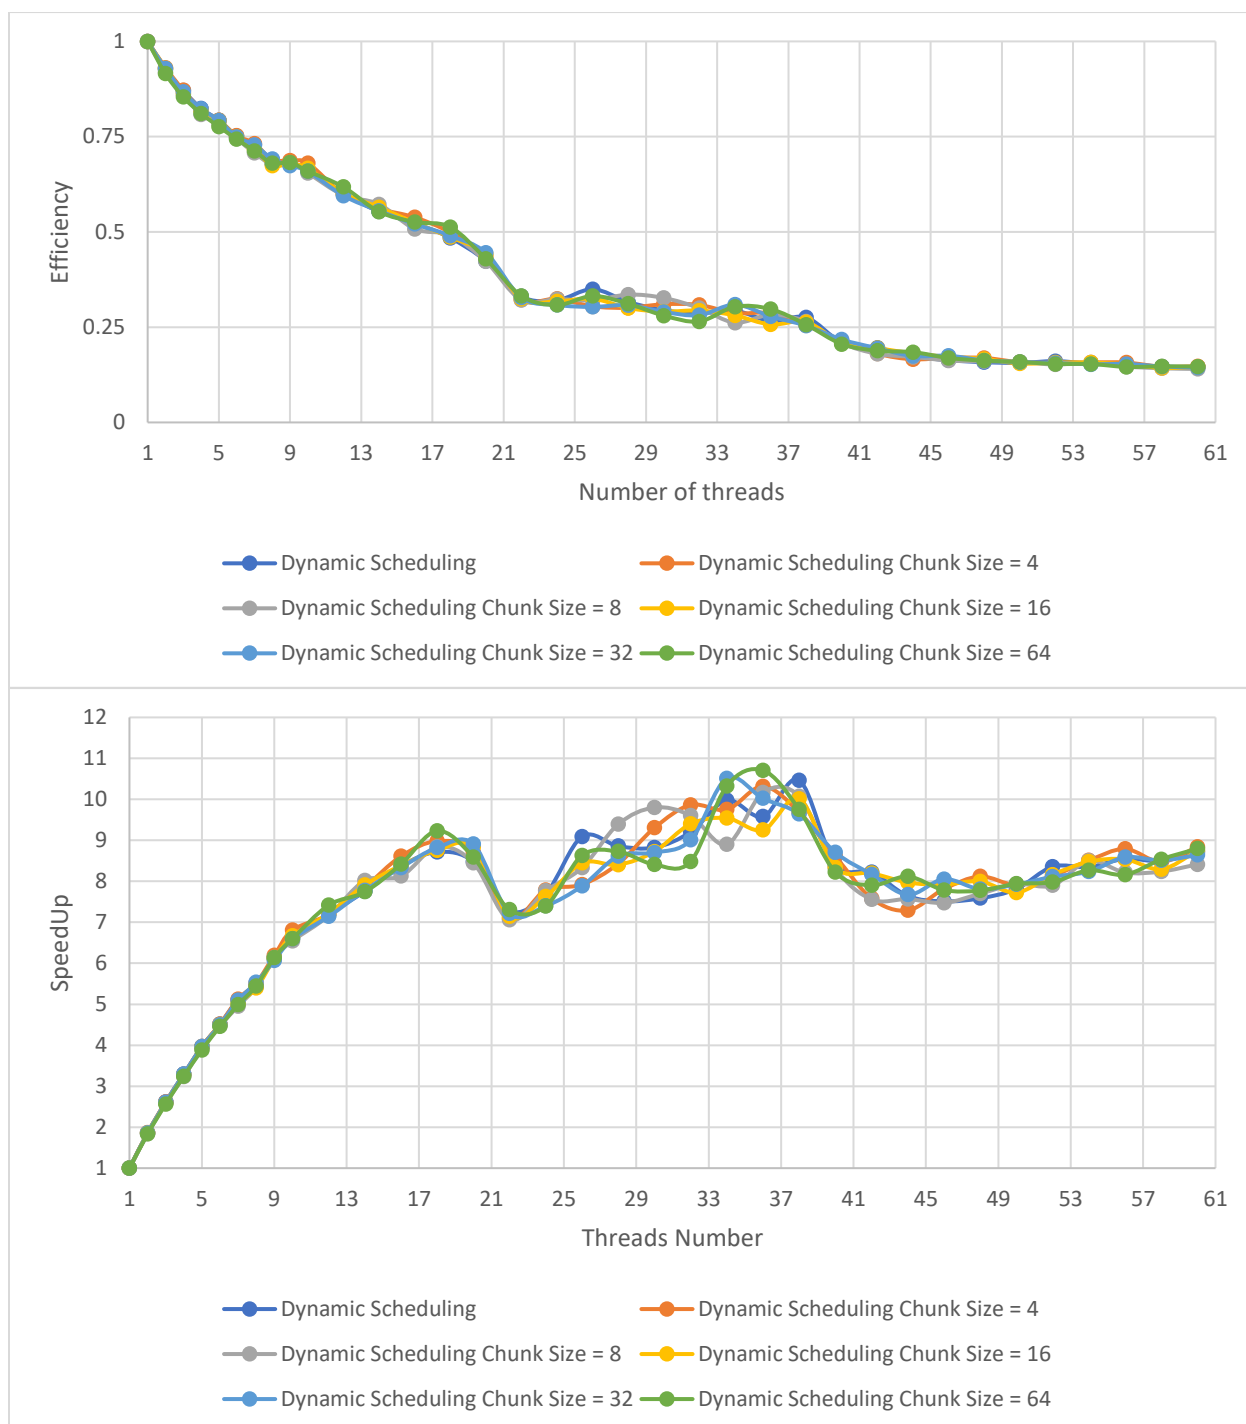
1. Static, Chunk Size = 4,8,16,32,64 Speedup + Efficiency



2. Guided, Chunk Size = 4,8,16,32,64 Speedup + Efficiency

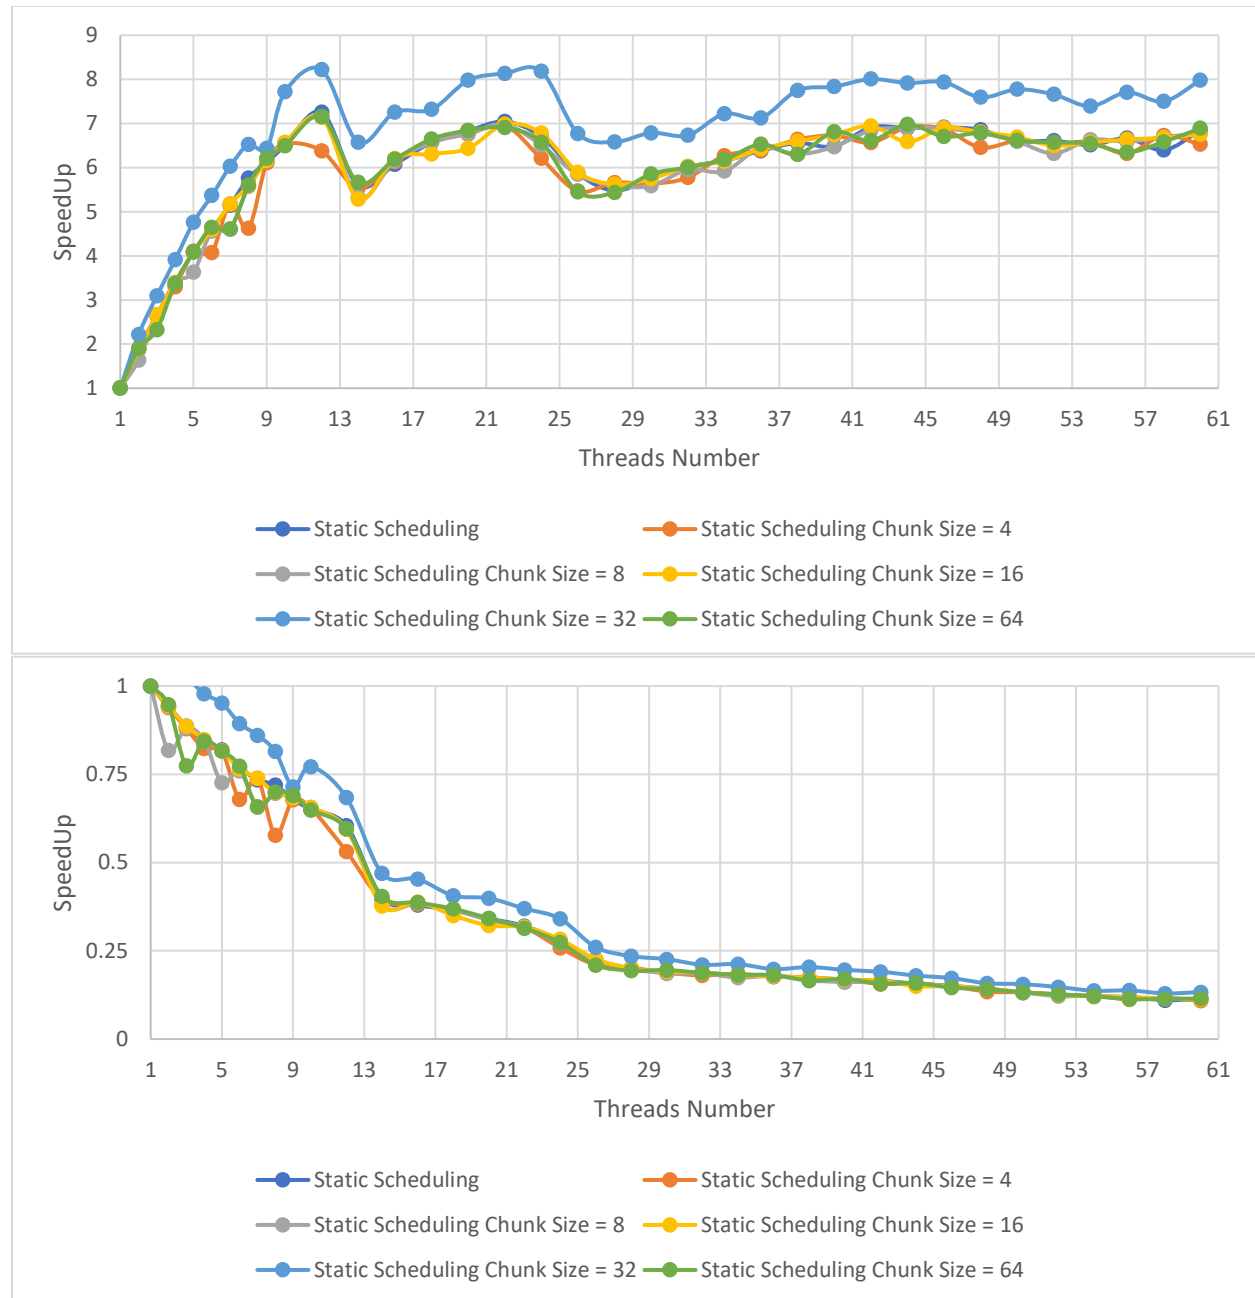


3. Dynamic, Chunk Size = 4,8,16,32,64 Speedup + Efficiency

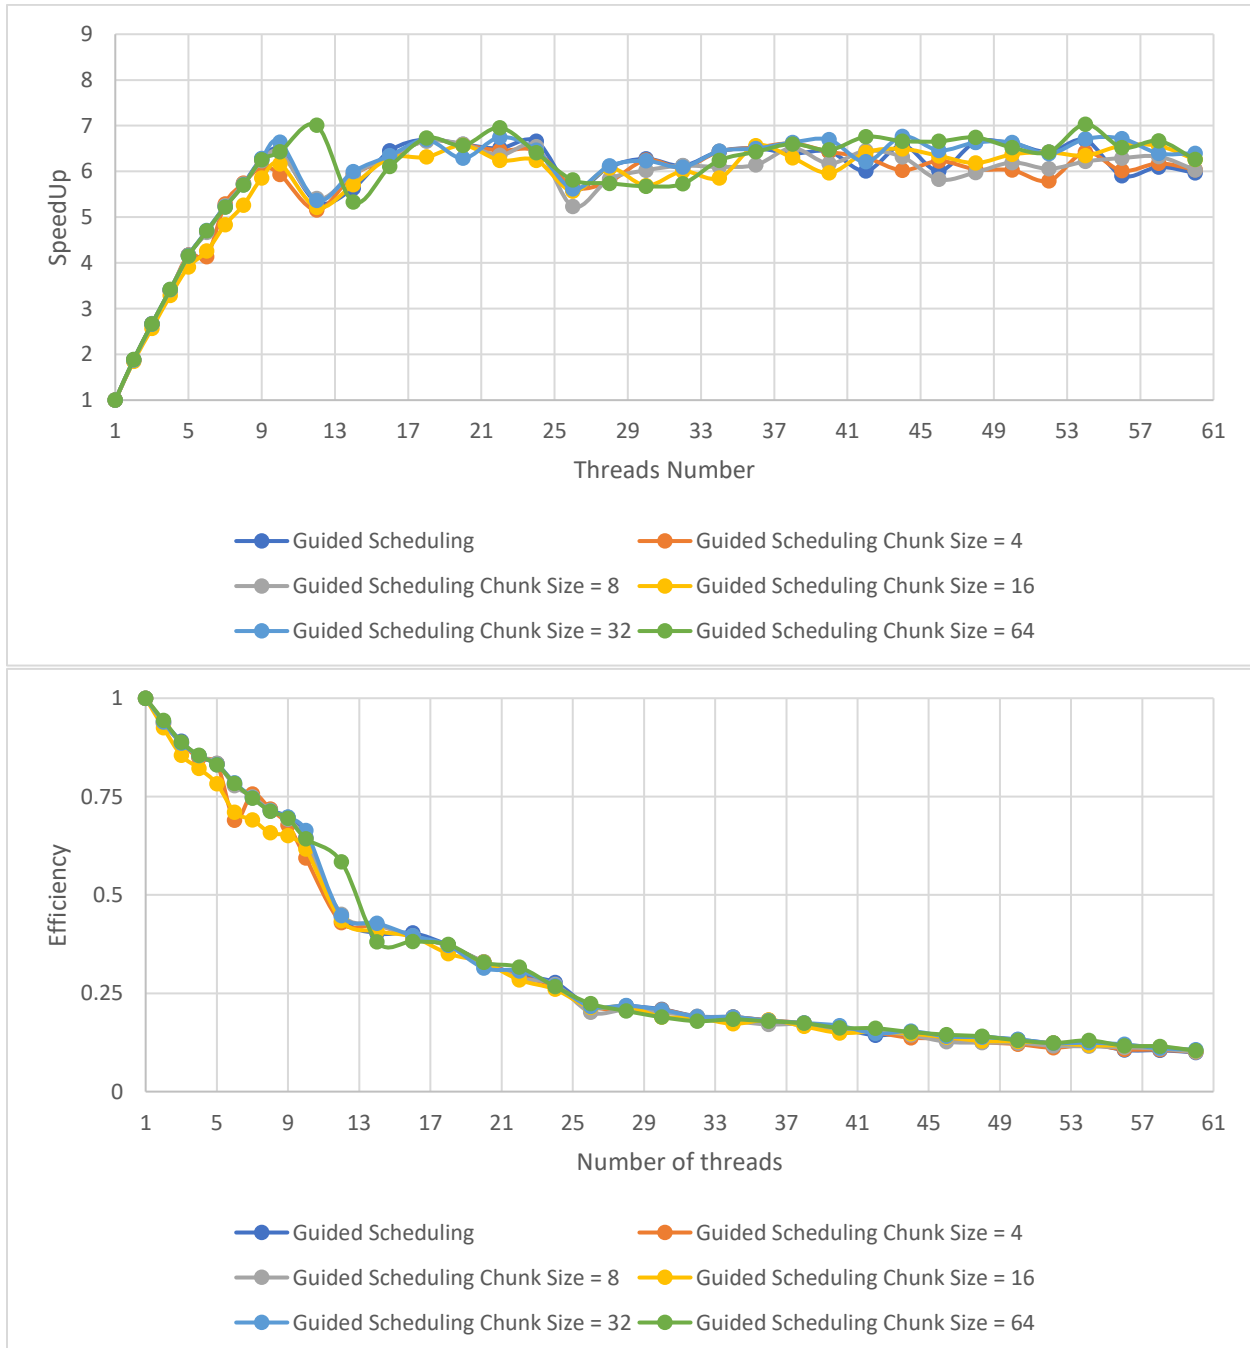


Running with -O3 on IBM cluster


1. Static, Chunk Size = 4,8,16,32,64 Speedup + Efficiency



2. Guided, Chunk Size = 4,8,16,32,64 Speedup + Efficiency



3. Dynamic, Chunk Size = 4,8,16,32,64 Speedup + Efficiency



Speed Up												
Scheduling \ Chunk-size	1		4		8		16		32		64	
	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM
static	11.20	7.254	10.45	6.954	9.823	7.153	10.23	7.139	10.26	8.216	10.32	7.151
	T:36	T:12	T:36	T:22	T:38	T:12	T:36	T:12	T:34	T:12	T:38	T:12
dynamic	10.46	6.929	10.31	6.914	10.16	7.134	10.01	7.083	10.51	7.042	10.70	7.057
	T:38	T:12	T:36	T:12	T:36	T:12	T:38	T:12	T:34	T:12	T:36	T:12
guided	9.688	6.704	10.39	6.679	10.06	6.634	9.981	6.572	10.32	6.745	10.56	7.032
	T:32	T:18	T:34	T:18	T:36	T:18	T:38	T:20	T:34	T:22	T:34	T:12

CPU - time												
Scheduling \ Chunk-size	1		4		8		16		32		64	
	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM	HPSL	IBM
static	13.848	20.405	14.831	21.255	15.742	20.748	15.201	20.861	15.134	21.118	14.968	20.798
	T:36	T:12	T:36	T:22	T:38	T:12	T:36	T:12	T:34	T:12	T:38	T:12
dynamic	14.89	21.352	15.13	21.508	15.167	20.772	15.519	21.04	14.785	21.039	14.394	21.010
	T:38	T:12	T:36	T:12	T:36	T:12	T:38	T:12	T:34	T:12	T:36	T:12
guided	16.059	22.125	14.975	22.227	15.287	22.351	15.611	22.705	15.135	21.920	14.657	21.199
	T:32	T:18	T:34	T:18	T:36	T:18	T:38	T:20	T:34	T:22	T:34	T:12

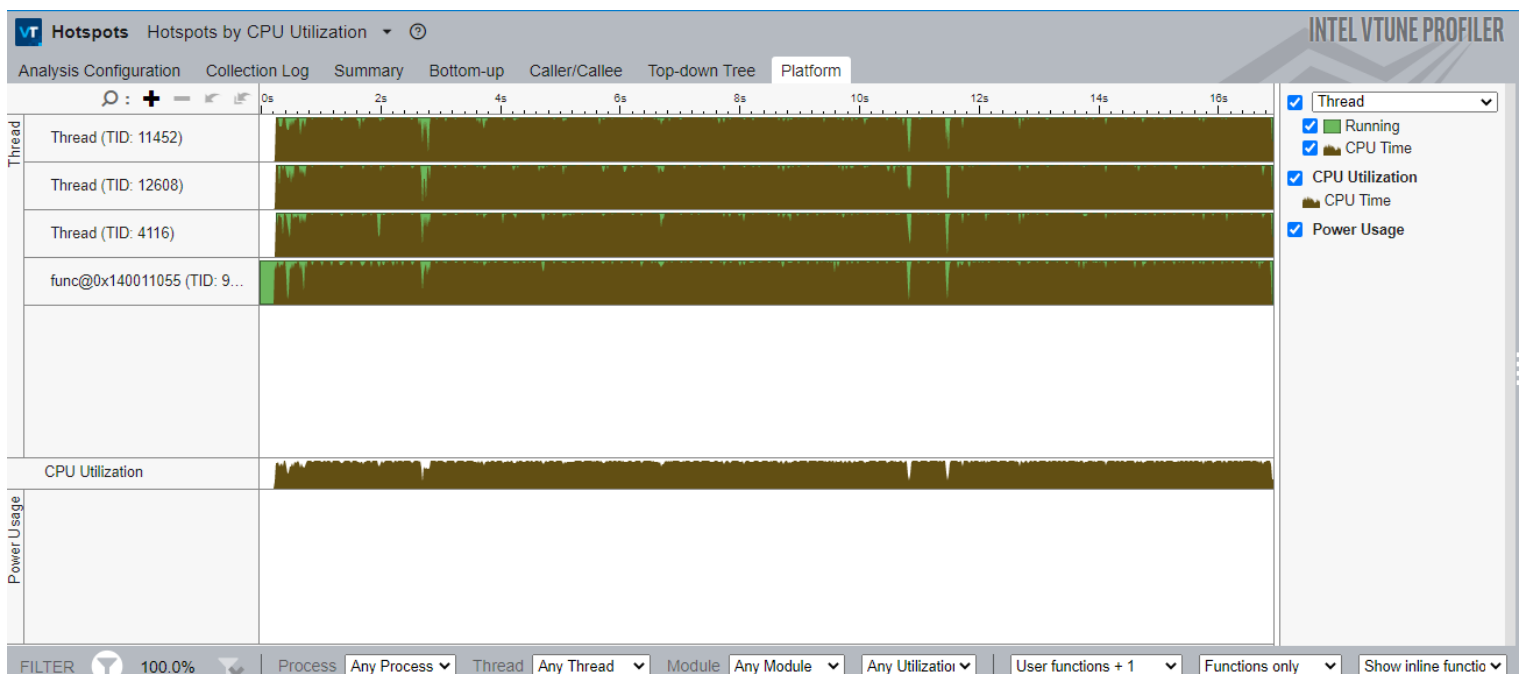

Parallel Algorithm (III) – implementation without opencv

I was unable to replace or create a function to write a movie in MJPEG format. Instead, I saved every second of the video as a bitmap image. For reading and writing BMP files I created a structure below for BMP header, copy BMP footer, check if file is a BMP file, read file, copy data and write data.

Hotspots by CPU Utilization

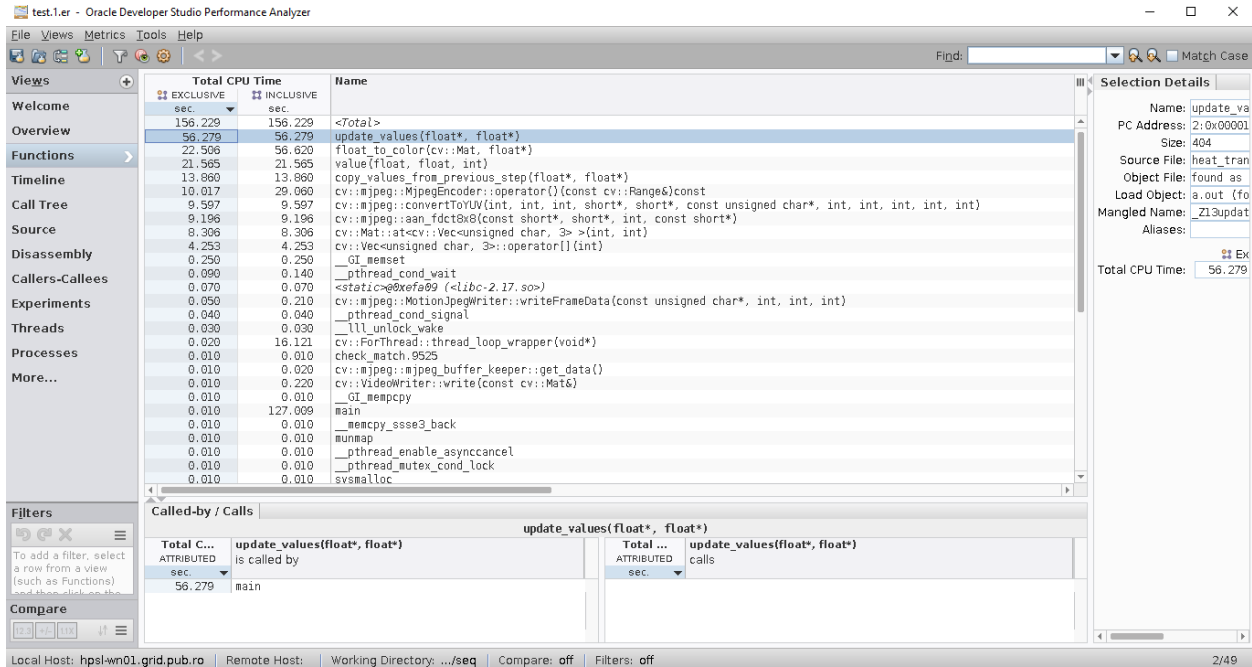
Elapsed Time: 16.909 s
Total Thread : 4
Paused Time : 0 s

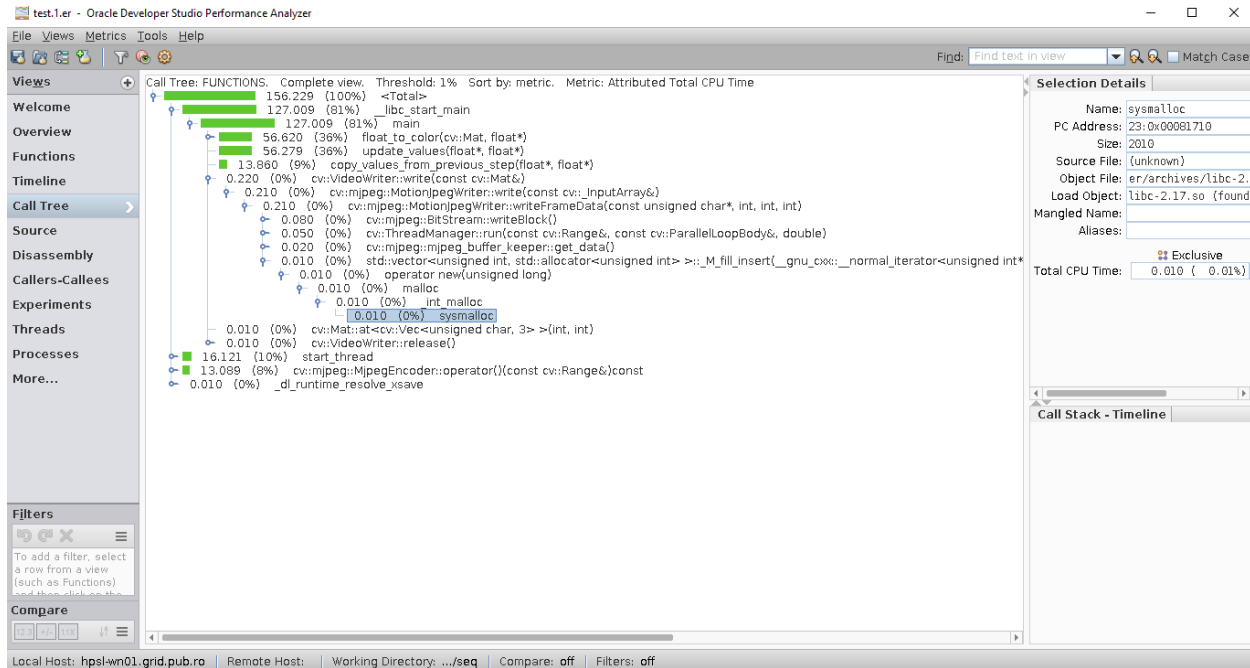
Effective CPU Utilization Histogram



Profiling with Solaris Studio


Profiling for Sequential Algorithm






Profiling for Sequential Algorithm II

Total CPU Time	Name
EXCLUSIVE	INCLUSIVE
sec.	sec.
167.797	167.797
64.055	64.055
25.017	64.595
20.855	20.855
14.720	14.720
13.860	13.860
10.597	24.157
6.775	6.775
6.555	6.555
4.863	4.863
0.230	0.230
0.080	0.080
0.060	0.210
0.040	0.040
0.030	0.030
0.020	0.020
0.010	0.020
0.010	0.010
0.010	143.610
0.010	0.010
0.	0.010
0.	13.429
0.	0.030
0.	0.220
0.	0.040
0.	0.010
0.	0.220
0.	0.030
0.	0.030
0.	0.030
0.	143.610
0.	13.429
0.	13.429
0.	0.010

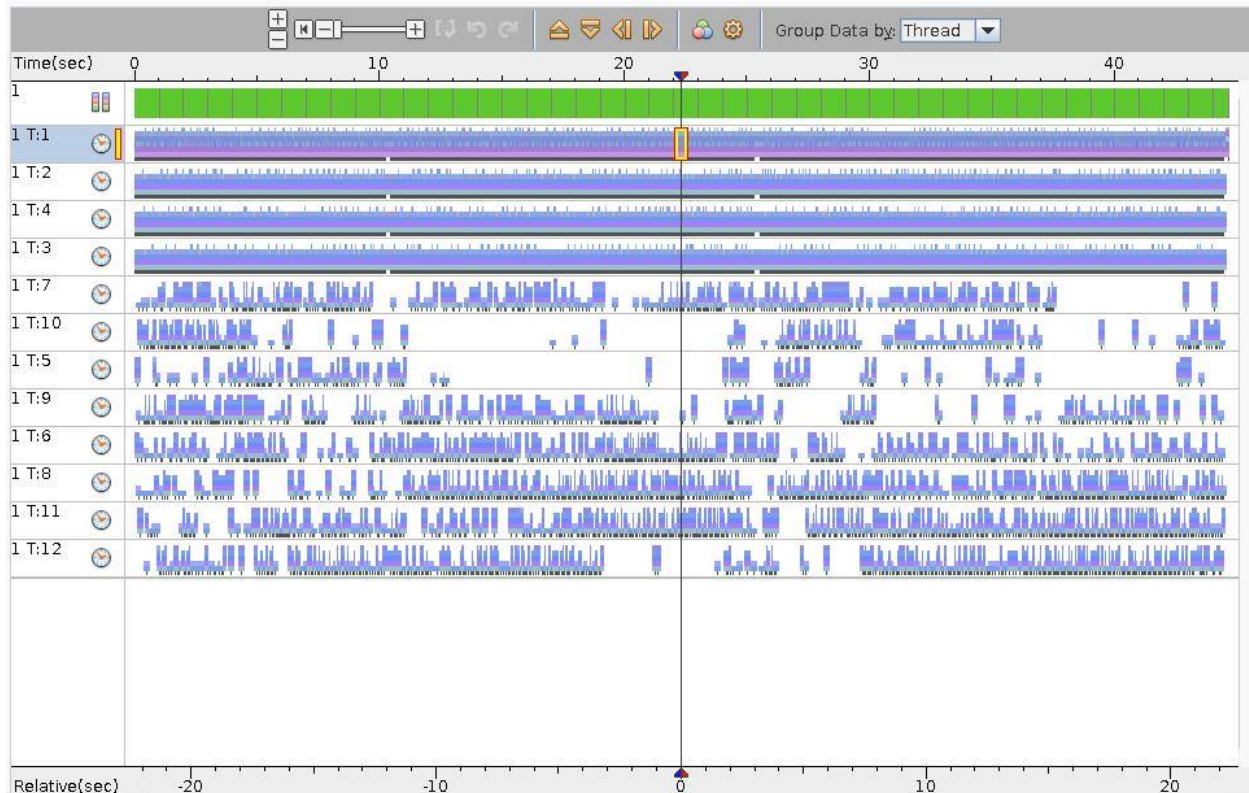


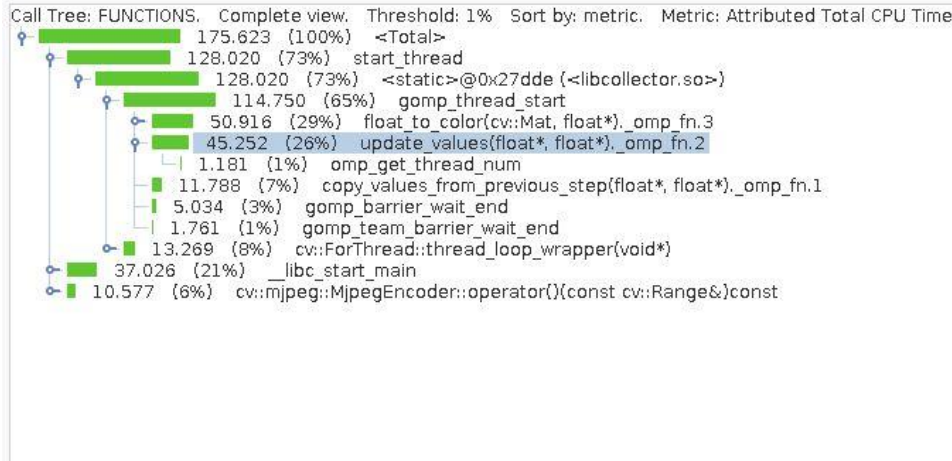
Call Tree: FUNCTIONS. Complete view. Threshold: 1% Sort by: metric. Metric: Attributed Total CPU Time



Profiling for Parallel Algorithm I

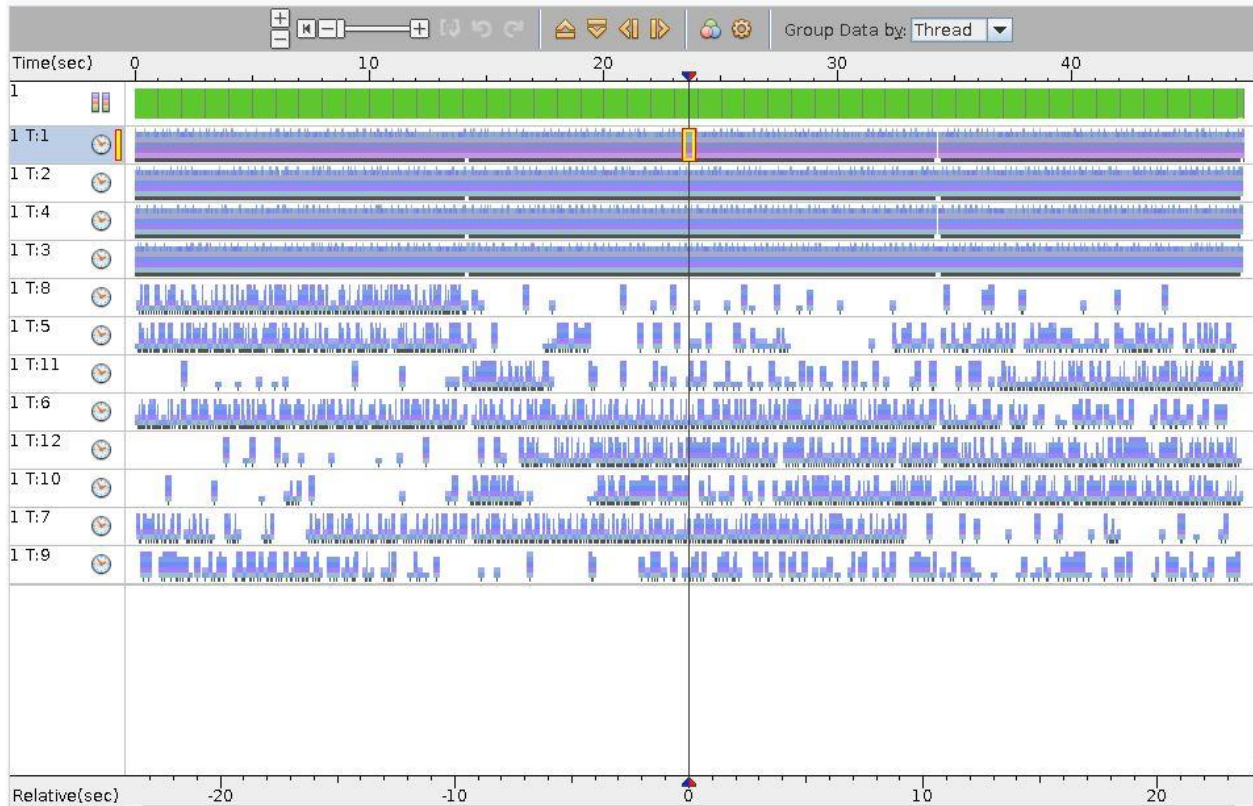
Total CPU Time		Name	Selection Details
EXCLUSIVE	INCLUSIVE		
sec.	sec.		
175.623	175.623	<Total>	Name: update_values(float*,
58.851	60.422	update_values(float*, float*)_omp_fn.2	PC Address: 2:0x00001C4B
29.371	67.377	float_to_color(cv::Mat, float*)_omp_fn.3	Size: 525
19.914	19.914	value(float, float, int)	Source File: t_transfer_omp1.cpp
15.381	15.381	copy_values_from_previous_step(float*, float*)_omp_fn.1	Object File: /a.out_my2NjyEF4n7)
13.399	13.399	cv::Mat::at<cv::Vec<unsigned char, 3>>(int, int)	Load Object: a.out (found as test..
8.956	23.747	cv::mjpeg::MjpegEncoder::operator()(const cv::Range&)/const	Mangled Name: _Z13update_valuesPi
7.485	7.485	cv::mjpeg::convertToYUV(int, int, short*, short*, const unsigned char*, int, int, int, int)	Aliases:
7.125	7.125	cv::mjpeg::aan_fdct8x8(const short*, short*, int, const short*)	
5.044	5.044	gomp_barrier_wait_end	
4.693	4.693	cv::Vec<unsigned char, 3>::operator[](int)	
3.292	3.292	gomp_team_barrier_wait_end	
1.571	1.571	omp_get_thread_num	
0.180	0.180	_GI_memset	
0.090	0.240	cv::mjpeg::MotionJpegWriter::writeFrameData(const unsigned char*, int, int, int)	
0.080	0.110	_pthread_cond_wait	
0.050	0.060	<static>@0x0a09 (<libc-2.17.so>)	
0.030	0.030	cv::mjpeg::mjpeg_buffer_keeper::get_data()	
0.020	0.020	_l1_unlock_wake	
0.020	0.020	_pthread_cond_signal	
0.010	0.010	_close_nocancel	
0.010	13.269	cv::ForThread::thread_loop_wrapper(void*)	
0.010	0.010	cv::mjpeg::BitStream::putBytes(const unsigned char*, int)	
0.010	0.010	cv::parallel_for_pthreads(const cv::Range&, const cv::ParallelLoopBody&, double)	
0.010	0.010	_libc_disable_asynccancel	
0.010	0.010	_pthread_disable_asynccancel	
0.010	0.010	std::vector<unsigned int, std::allocator<unsigned int>> >::M_fill_insert(__gnu_cxx::__normal_iterator<unsigned int*, s	
0.	4.143	copy_values_from_previous_step(float*, float*)	
0.	0.010	cv::detail::PtrOwnerImpl<cv::mjpeg::MotionJpegWriter, cv::DefaultDeleter<cv::mjpeg::MotionJpegWriter>> >::deleteSelf	
0.	0.050	cv::mjpeg::BitStream::writeBlock()	
0.	0.250	cv::mjpeg::MotionJpegWriter::write(const cv::InputArray&)	
0.	0.040	cv::ThreadManager::run(const cv::Range&, const cv::ParallelLoopBody&, double)	
0.	0.010	cv::VideoWriter::release()	
0.	0.250	cv::VideoWriter::write(const cv::Mat&)	
0.	17.272	float_to_color(cv::Mat, float*)	
0.	0.010	fseek	
0.	0.050	fwrite	






Profiling for Parallel Algorithm II

Total CPU Time		Name
EXCLUSIVE sec.	INCLUSIVE sec.	
185.780	185.780	<Total>
58.141	59.682	update_values(float*, float*)
31.992	72.321	float_to_color(cv::Mat, float*)
21.065	22.616	copy_values_from_previous_step(float*, float*)
20.925	20.925	value(float, float, int)
13.870	13.870	cv::Mat::at<cv::Vec<unsigned char, 3>>(int, int)
9.246	24.397	cv::mjpeg::MjpegEncoder::operator()(const cv::Range&)const
7.805	7.805	cv::mjpeg::convertToYUV(int, int, int, short*, short*, const unsigned char*, int, int, int, int, int)
7.165	7.165	cv::mjpeg::aan_fdct8x8(const short*, short*, int, const short*)
6.585	6.585	gomp_team_barrier_wait_end
4.933	4.933	cv::Vec<unsigned char, 3>::operator[](int)
3.693	3.693	omp_get_thread_num
0.180	0.180	_GI_memset
0.080	0.130	cv::mjpeg::MotionJpegWriter::writeFrameData(const unsigned char*, int, int, int)
0.020	0.020	cv::mjpeg::mjpeg_buffer_keeper::get_data()
0.020	0.020	__pthread_cond_wait
0.010	0.010	__close_nocancel
0.010	0.010	fwrite
0.010	161.343	main_omp_fn.1
0.010	0.010	pthread_mutex_unlock
0.010	0.010	<static>@0x7769 {<libc-2.17.so>}
0.010	0.010	std::vector<unsigned int, std::allocator<unsigned int>>::_M_fill_insert(_gnu_cxx::__normal_iterator<unsigned int*, s
0.	0.010	cv::createMotionJpegWriter(const cv::String&, double, cv::Size<int>, bool)
0.	0.010	cv::detail::PtrOwnerImpl<cv::mjpeg::MotionJpegWriter, cv::DefaultDeleter<cv::mjpeg::MotionJpegWriter>>::deleteSelf()
0.	13.890	cv::ForThread::thread_loop_wrapper(void*)
0.	0.010	cv::makePtr<cv::mjpeg::MotionJpegWriter, cv::String, double, cv::Size<int>, bool>(const cv::String&, const double&,
0.	0.010	cv::mjpeg::BitStream::writeBlock()
0.	0.130	cv::mjpeg::MotionJpegWriter::write(const cv::InputArray&)
0.	0.010	cv::ThreadManager::run(const cv::Range&, const cv::ParallelLoopBody&, double)
0.	0.010	cv::VideoWriter::open(const cv::String&, int, double, cv::Size<int>, bool)
0.	0.010	cv::VideoWriter::release()
0.	0.010	cv::VideoWriter::VideoWriter(const cv::String&, int, double, cv::Size<int>, bool)
0.	0.130	cv::VideoWriter::write(const cv::Mat&)
0.	0.010	_fopen_internal
0.	40.068	GOMP_parallel
0.	121.275	gomp_thread_start
0.	0.010	_IO_new_file_close_it




Call Tree: FUNCTIONS. Complete view. Threshold: 1% Sort by: metric. Metric: Attributed Total CPU Time



Profiling for Sequential Program without OpenCV library

Total CPU Time		Name
EXCLUSIVE sec.	INCLUSIVE sec.	
135.625	135.625	<Total>
64.775	64.775	update_values(float*, float*)
25.748	52.987	float_to_color(std::vector<unsigned char, std::allocator<unsigned char> >&, float*)
23.216	23.216	value(float, float, int)
14.670	14.670	copy_values_from_previous_step(float*, float*)
4.033	4.033	std::vector<unsigned char, std::allocator<unsigned char> >::operator[](unsigned long)
1.761	1.761	writetv
0.801	0.801	_open_nocancel
0.590	0.590	_close_nocancel
0.010	0.010	_IO_link_in
0.010	0.010	operator new(unsigned long)
0.010	0.010	std::ios_base::ios_base()
0.	3.182	BMP::writeBMP(const char*)
0.	1.761	BMP::write_headers_and_data(std::basic_ofstream<char, std::char_traits<char> >&)
0.	0.811	_fopen_internal
0.	0.590	_IO_new_file_close_it
0.	0.801	_IO_new_file_fopen
0.	0.010	_IO_new_file_init_internal
0.	135.625	_libc_start_main
0.	135.625	main
0.	0.590	_new_fclose
0.	0.590	std::basic_filebuf<char, std::char_traits<char> >::close()
0.	0.010	std::basic_filebuf<char, std::char_traits<char> >::_M_allocate_internal_buffer()
0.	0.821	std::basic_filebuf<char, std::char_traits<char> >::open(const char*, std::_ios_Openmode)
0.	1.761	std::basic_filebuf<char, std::char_traits<char> >::xspn(const char*, long)
0.	0.590	std::_basic_file<char>::close()
0.	0.811	std::_basic_file<char>::open(const char*, std::_ios_Openmode, int)
0.	1.761	std::_basic_file<char>::xspn_2(const char*, long, const char*, long)
0.	0.590	std::basic_ofstream<char, std::char_traits<char> >::~basic_ofstream()
0.	0.831	std::basic_ofstream<char, std::char_traits<char> >::basic_ofstream(const char*, std::_ios_Openmode)
0.	1.761	std::basic_ofstream<char, std::char_traits<char> >::write(const char*, long)



Call Tree: FUNCTIONS. Complete view. Threshold: 1% Sort by: metric. Metric: Attributed Total CPU Time

- 135.625 (100%) <Total>
- 135.625 (100%) _libc_start_main
 - 135.625 (100%) main
 - 64.775 (48%) update_values(float*, float*)
 - 52.987 (39%) float_to_color(std::vector<unsigned char, std::allocator<unsigned char> >&, float*)
 - 14.670 (11%) copy_values_from_previous_step(float*, float*)
 - 3.182 (2%) BMP::writeBMP(const char*)
 - 0.010 (0%) std::vector<unsigned char, std::allocator<unsigned char> >::operator[](unsigned long)